Arab Republic of Egypt
Ministry of Higher Education
and Scientific Research

H.I.T Academy
Department of Computer Science
Graduation Project

THE EGYPTIAN AWARENESS LIBRARY

المكتبة المصرية للوعي

Three platforms — for verifying truth, preparing for emergencies, and understanding the mind.

SOFTWARE PROJECT DOCUMENTATION

Mission · Methodology · Requirements · Analysis · Design · Implementation · Testing

① **Mawthooq** موثوق — Verification

② **Mosta'ed** مستعد — Preparedness

③ **Motazen** متزن — Psychological Literacy

Document type: Graduation Project — Technical Documentation

Subject: The Egyptian Awareness Library and its three platforms

Version: 1.0 · **Date:** July 2026 · **Language:** English

Status: Mawthooq deployed (live); Mosta'ed verified (pre-deploy); Motazen MVP complete

UNDER THE SUPERVISION OF

Dr. Bosy Mohamed

Eng. Eman Abdelwahed

PREPARED BY

Khaled Sayed Hamed (Team Lead) · Mohamed Sherief Mohamed

Aya Mohamed Abdullah El-Sebki · Donia Ali Sharaby

Ziad Fateen Kamal · Yasmine Mohamed El-Araby

Declaration

This document is the technical documentation of the **Egyptian Awareness Library** and its three platforms — **Mawthooq**, **Mosta'ed**, and **Motazen**. We declare that the work described herein is our own, that all external material — including the academic sources that ground the platforms' content and the documentation-methodology references used in Chapter 2 — is acknowledged and cited, and that every quantitative figure reported was measured directly from the live systems rather than estimated. Any statement in the platforms' content that could not be verified against a real, fetched source was, by design, excluded (see "The One Law," §3.2).

Project title: **The Egyptian Awareness Library**	Date: July 2026
Signature:	Supervisor: ....

Acknowledgement

We are grateful to our supervisor and to the faculty whose guidance shaped this project, and to the peers and reviewers who tested the platforms and challenged their claims. We acknowledge, in particular, a debt to the principle that gives the whole library its integrity — that nothing should be told to a person that cannot be traced to a real source — and to the researchers and public-health bodies whose openly available work made a source-verified platform possible.

This project is dedicated to every Arabic-speaking reader who has ever been handed a confident falsehood about the world, about an emergency, or about their own mind — and deserved better.

Abstract

The **Egyptian Awareness Library** (المكتبة المصرية للوعي, "EAL") is a family of Arabic-first digital platforms built to raise public awareness and defend Arabic-speaking users against a polluted information environment. Its members share a single, uncompromising engineering principle — **The One Law**: *no factual content claim reaches a user unless it is extracted from a real, fetched, hash-verified source, and every stored quotation appears verbatim in that source.* (The law governs the platforms' teachable content; ordinary interface text and navigation are outside its scope.) Artificial intelligence is used only as an *assembler* of grounded content, never as a *factory* that invents it. This law is enforced mechanically by automated **firewall validators** that fail the build if any citation cannot be re-verified.

The library comprises three platforms, documented here in order. **Mawthooq** (موثوق, "trusted") is the flagship **verification** platform — a large Next.js/React web application of roughly **153 pages and 95 API endpoints** providing claim-debunking, media-forensics, bias- and fallacy-detection, and a critical-thinking curriculum; it is deployed and live. **Mosta'ed** (مستعد, "prepared") is a **preparedness** platform delivering source-verified emergency and first-aid guidance localised to Egypt, with every card validated against fetched, hash-stamped public-health sources. **Motazen** (متزن, "balanced") is a **psychological-literacy** platform of **118 grounded claims** across a six-phase curriculum, teaching users *how to think about* mental health with full epistemic honesty.

This document specifies the library and its platforms as a software-engineering report. It defines documentation quality itself; reviews the problem domain and the shared methodology; states functional and non-functional requirements; presents the four core UML views — **Use Case, Activity, Sequence,** and **Class** diagrams — of the shared system; details the architecture, the common content data-model, and the three visual identities; describes the implementation of each platform and its verification pipeline; reports testing results; and closes with limitations and a roadmap. Every quantitative fact herein was measured directly from the live systems, in keeping with the library's own law.

Keywords: misinformation defense; content provenance; anti-fabrication; mental-health literacy; verification; preparedness; digital health; Arabic-first design; UML; Egypt.

Table of Contents

Declaration 3

Acknowledgement 4

Abstract 5

List of Figures 9

List of Tables 10

PART I — PRELIMINARIES

1 Introduction 11

1.1 Background and Motivation 11

1.2 Problem Statement 11

1.3 Aim and Objectives 12

1.4 Scope 12

1.5 Audience, How to Read, and Document vs. PRD 12

1.6 Development Methodology and Process Model 13

2 Documentation Standards and Quality 14

2.1 What Documentation Is 14

2.2 Bad vs. Good Documentation 14

2.3 Maturity: Bare-Minimum to Ultra 15

2.4 The Standard Shape; Standards Adopted 16

PART II — THE SYSTEM AND ITS PLATFORMS

3 The Egyptian Awareness Library (System Overview) 18

3.1 Mission and Vision 18

3.2 The One Law (Anti-Fabrication Doctrine) 18

3.3 The Verification Doctrine 18

3.4 The Shared Content Pipeline 19

3.5 Brand Identity: The Hall of Truth 19

3.6 The Family of Platforms 19

3.7 Source Tiers, Licensing, and the Content Lifecycle 21

4 Platform I — Mawthooq (Verification) 22

4.1 Purpose and Users 22

4.2 Feature Set 22

4.3 Architecture and Access Control 24

4.4 Status ... 25
4.5 The Engine Pattern and the Calibrated Verdict ... 25
4.6 The Critical-Thinking Curriculum ... 25

5 Platform II — Mosta'ed (Preparedness) ... 26
5.1 Purpose and Users ... 26
5.2 The Card Provenance Pipeline ... 26
5.3 Egyptian Localisation (Gate 6) ... 27
5.4 Status ... 27

6 Platform III — Motazen (Psychological Literacy) ... 28
6.1 Purpose and Users ... 28
6.2 The Six-Phase Curriculum ... 28
6.3 The Nine Surfaces ... 29
6.4 Safety and the Four Honesty States ... 29
6.5 The Six Domains in Detail ... 29
6.6 Example Claims (as Published) ... 30

PART III — ANALYSIS, DESIGN, AND DELIVERY

7 System Analysis and Requirements ... 32
7.1 Use Case Diagram and Specifications ... 32
7.2 Activity Diagrams ... 34
7.3 Sequence Diagrams ... 36
7.4 Claim Lifecycle (State Machine) ... 37
7.5 Functional and Non-Functional Requirements ... 37

8 System Design ... 39
8.1 Layered Architecture ... 39
8.2 Class Diagram and Data Model ... 39
8.3 Visual Identity and UX ... 41
8.4 Deployment and Maintenance View ... 41

9 Implementation ... 43
9.1 Technology Stack ... 43
9.2 The Content Pipeline ... 43
9.3 Key Algorithms ... 43
9.4 Repository and Build Structure ... 44

10 Testing and Verification ... 46
10.1 The Firewall as an Automated Test ... 46

10.2 Verification Results 46
10.3 Browser and Safety Verification 46
10.4 Ethics, Safety, and Risk Analysis 47

11 Results and Evaluation 48
11.1 Outcomes Against Objectives 48
11.2 Development Process and Lessons Learned 49
11.3 Competitive Differentiation 49

12 Conclusion and Future Work 51
12.1 Summary 51
12.2 Limitations 51
12.3 Future Work 51

BACK MATTER

References 53

Appendix A — Requirements Traceability Matrix 54

Appendix B — The Claim/Card Object Schema 56

Appendix C — Glossary and Abbreviations 57

Appendix D — The Six Quality Strategies 58

Appendix E — Content Inventory 59

Appendix F — Complete Platform Inventory (all 153 routes + surfaces) 60

Appendix G — Team Contributions and Division of Labour 63

List of Figures

Figure	Title	Page
Figure 1.1	The Iterative, Content-First Process Model	13
Figure 3.1	The Shared Content Pipeline (conceptual)	19
Figure 3.2	The Family of Platforms	20
Figure 7.1	Use Case Diagram — The Egyptian Awareness Library	32
Figure 7.2	Activity Diagram — The One-Law Content Pipeline	35
Figure 7.3	Activity Diagram — The Learner's Journey (Motazen)	35
Figure 7.4	Sequence Diagram — Assistant Query with Crisis Routing	36
Figure 7.5	Sequence Diagram — Content Verification (the Firewall)	37
Figure 7.6	State Machine — Claim Lifecycle	37
Figure 8.1	Layered Architecture	39
Figure 8.2	Class Diagram — Shared Content Model	40
Figure 8.3	Deployment Diagram	41
Figure 11.1	Content Library in Numbers (honesty states; domains)	48

List of Tables

Table	Title	Page
Table 1.1	The Problem, in Numbers	11
Table 1.2	Reading Map by Audience	13
Table 2.1	Concrete Anti-Patterns and Their Corrections	15
Table 2.2	Documentation Maturity Spectrum	16
Table 2.3	Global Documentation Standards and How Applied	17
Table 3.1	The Three Platforms at a Glance	21
Table 3.2	Source Tiers	21
Table 3.3	Positioning Against Typical Alternatives	20
Table 4.1	Mawthooq — Representative Feature Modules	23
Table 4.2	Mawthooq — Flagship Capabilities	24
Table 5.1	Anatomy of a Readiness Card	27
Table 6.1	Phase-1 Domain Inventory	30
Table 7.1	Actors and Responsibilities	33
Table 7.2–7.4	Use Case Specifications (Assistant, Verify, Author)	33
Table 7.5	Functional Requirements	38
Table 7.6	Non-Functional Requirements	38
Table 8.1	The Claim/Card Object — Field Dictionary	40
Table 9.1	Technology Stack by Platform	43
Table 10.1	Risks and Mitigations	47
Table 11.1	Library Metrics Across Platforms	48
Table 11.2	Key Lessons from the Build	49
Table 11.3	Capability Comparison Against Rival Categories	50
Table 12.1	Four Pathways to Reach Scale	52
Table A.1–A.2	Requirements Traceability (Functional, Non-Functional)	54
Table E.1–E.3	Content Inventory (Phases, Distribution, Surfaces)	59
Table F.1–F.2	Complete Platform Inventory (Mawthooq routes; Motazen surfaces)	60

1 Introduction

1.1 Background and Motivation

The modern information environment is polluted. Confident, credentialed-sounding claims — miracle cures, one-line neuro-explanations, manipulated media, and "detox" scams — spread precisely because they are simple and certain, and they spread fastest in languages and regions least served by trustworthy, accessible alternatives. Arabic-speaking users in Egypt and the wider Arab world face this problem acutely across three high-stakes arenas: **public claims** they cannot easily verify, **emergencies** they are unprepared for, and their own **mental health**, about which the surrounding conversation is often stigmatising and wrong.

The **Egyptian Awareness Library** was created to meet all three. It is not a single app but a **family** of platforms unified by one methodology and one mission: to give ordinary people the tools to tell what is true, to act when it matters, and to think about themselves with clarity. Where a typical "awareness" project relies on an author's authority, the EAL relies on **verifiable provenance** — every claim it makes can be traced back to a real, fetched source, and the system is engineered so that it *cannot* say otherwise.

The scale of the underlying problem, drawn from the library's own sourced content, motivates the work (Table 1.1).

Fact	Significance
Mental disorders are among the **top-ten** causes of global disease burden, with **no decline 1990–2019**.	The status quo is not working; awareness and literacy are part of what is missing.
The **treatment gap exceeds 70%** in many low- and middle-income settings.	Most people who could benefit never reach care — first, a literacy and access failure.
About **12% of humanity supplies over 90%** of behavioural-science research subjects (the "WEIRD" problem).	Imported findings must be flagged where they were not tested on the user's own population.
In a landmark replication project, only about **36% of 100** psychology studies replicated.	Even "known" science is provisional; honesty about uncertainty is a design requirement, not a caveat.

Table 1.1 — The problem, in numbers (from sources cited within the platforms).

1.2 Problem Statement

Stated precisely: *Arabic-speaking people lack trustworthy, culturally-situated, and epistemically-honest digital tools for verifying public claims, preparing for emergencies, and understanding mental health.* Every candidate solution faces the same hard sub-problem: **how to stop an automated content system from asserting plausible falsehoods.** A platform that lets a language model invent reassuring or authoritative-sounding statements is not an awareness tool; it is a misinformation vector with a trustworthylooking interface. The core engineering problem of the EAL is therefore as much about **what the system must refuse to publish** as about what it publishes.

1.3 Aim and Objectives

The aim is to design, implement, and verify a family of awareness platforms whose every user-facing *factual claim* is provably grounded in a real source. The objectives are:

- **O1 — A shared anti-fabrication engine.** Engineer a common "One Law" pipeline and automated firewall reused across all platforms.
- **O2 — Mawthooq:** deliver a comprehensive verification and critical-thinking platform (claim-checking, media forensics, bias/fallacy detection, curriculum).
- **O3 — Mosta'ed:** deliver source-verified, Egypt-localised emergency and preparedness guidance.
- **O4 — Motazen:** deliver an epistemically-honest psychological-literacy curriculum with multi-modal, safety-bounded surfaces.
- **O5 — Cultural situation and honesty:** author Arabic-first content, respect cultural and religious meaning-making, classify every claim by honesty state, and flag where evidence was not tested on the user's own population.

1.4 Scope

In scope. Three Arabic-first web platforms; a shared source-verification pipeline and firewall; the content libraries of each; and the interactive surfaces through which they are delivered. **Out of scope (hard non-goals):** none of the platforms is a substitute for a professional — Mawthooq does not render legal or medical verdicts of record; Mosta'ed is not an emergency dispatcher; Motazen is not a diagnostic tool, therapy, or crisis service. Each *detects* and *routes* to real human help rather than replacing it. These boundaries are treated as testable requirements (see §7, §10).

1.5 Audience, How to Read, and Document vs. PRD

This report serves examiners, engineers, and content reviewers. Examiners should read Chapters 1–3 and 10–12; engineers, Chapters 7–9 and the appendices; reviewers, Chapters 3–6. Following the Diátaxis principle that different needs demand different writing, each chapter holds a single purpose. Finally, this document is distinct from the individual product PRDs (such as the Arabic `Motazen_PRD_v1.md`): the PRDs argue *what* and *why* for one product; this document is the **engineering documentation** of the whole library — an SRS-plus-design-plus-implementation report, in English, with the formal UML views required for academic submission.

Reader	Read first	Why
Examiner / supervisor	Ch 1–3, 10–12	The argument, the methodology, and the evidence and limitations.
Engineer (extending it)	Ch 7–9, Appendices A–B, E	The models, the data schema, the pipeline, and the traceability.
Content author / clinician	Ch 3, §6.4–6.6, §10.4	The One Law, the honesty states, real example content, and the safety analysis.
Product stakeholder	Abstract, Ch 4–6, §12.3	What each platform does and the pathways to reach.

Table 1.2 — Reading map by audience.

1.6 Development Methodology and Process Model

The library was not built with a linear (Waterfall) process, and deliberately so. A Waterfall model assumes requirements are fixed up front; here they were not — the content of a literacy platform is discovered as its sources are read, and the central requirement (that *nothing unverifiable ships*) demands continuous checking rather than a single end-of-project test phase. The project therefore used an **iterative and incremental, content-first** process (Figure 1.1): content is produced in small increments, each of which passes through the full cycle — research and source, extract and assemble, verify (the firewall), classify and localise, build or refine the surface, and review — before the next increment begins. The firewall is embedded *inside* every iteration, not appended at the end, so verification is continuous and defects are caught at the increment that introduced them.

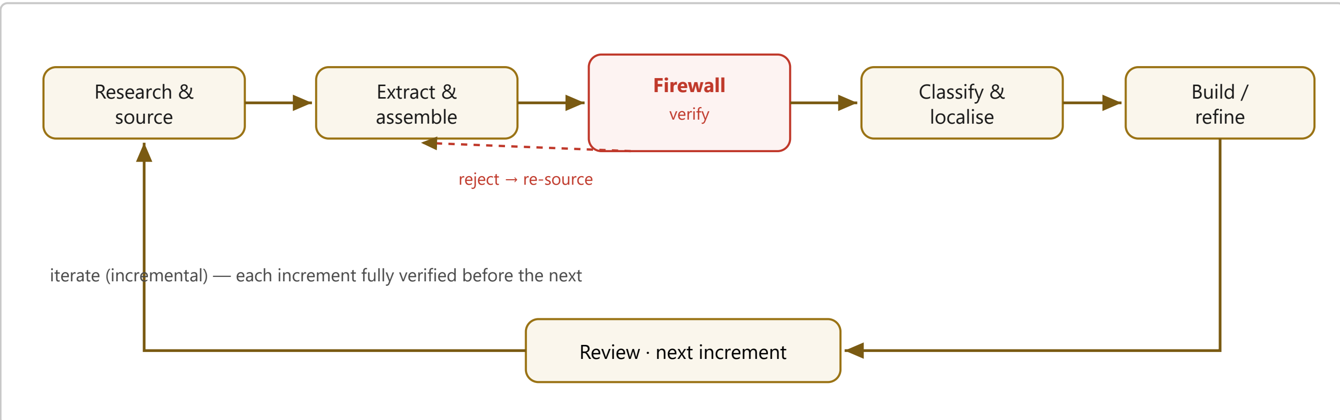

Figure 1.1 — The iterative, content-first process model. Verification (the firewall) is embedded in every increment, not deferred to a final phase.

This choice also explains a property visible throughout the document: because each increment is self-contained and durably recorded, the work is resilient to interruption, and the whole library can be regenerated and re-verified on demand (§9.4). Methodologically, the process is closest to an **incremental build with continuous verification** — Agile in its iteration and feedback, but with an unusually strict, automated quality gate that no increment may skip.

2 Documentation Standards and Quality

Because the quality of this deliverable is itself under assessment, this chapter defines what documentation *is*, separates *bad* from *good* using the technical-writing literature, sets out a maturity spectrum from bare-minimum to "ultra," and names the standard report shape and the standards this document follows. The structure adopted in this report was modelled on established academic and industry examples — university final-year-project report templates and a widely-used software-design-document template — and then extended.

2.1 What Documentation Is

Documentation is the durable, communicable record of a system: what it does, why it exists, how it is built, and how to use, extend, and trust it. It is the interface between a system and every human who is not its original author — including that author, months later. Good documentation is a first-class engineering artefact, not an afterthought; its job is to reduce the cost of understanding.

The **Diátaxis framework** (Procida) observes that technical documentation serves four different needs, each written differently: **tutorials** (learning-oriented lessons), **how-to guides** (goal-oriented directions), **reference** (accurate technical facts, free of interpretation), and **explanation** (understanding-oriented background answering *why*). This report is chiefly *reference* and *explanation*; the platforms themselves embody *tutorial* (their curricula) and *how-to* (their guidance surfaces).

2.2 Bad vs. Good Documentation

These definitions are drawn from the technical-writing literature, not from opinion. Practitioner surveys and technical-writing references converge on a consistent picture of how documentation fails and what distinguishes the good [3][4][5][6]. **Bad documentation** exhibits recognisable defects: **missing or incomplete content** — no coverage of essential tasks, features, or failure modes, and truncated examples that leave the reader stuck; **ambiguous or inconsistent language** — vague statements, undefined terms, contradictory instructions, and inconsistent naming; **poor structure and discoverability** — no clear headings, table of contents, or search-friendly layout, with important information buried or scattered; **drift and inaccuracy** — docs that no longer reflect the current system (worse than none, because they mislead with authority) and broken links; and **audience blindness and jargon** — text written for the author, dense with unexplained terminology [3][4][5]. The Diátaxis framework adds a structural failure: **mixing the four types** within one section, "the most common cause of confusing documentation" [1].

Good documentation, by contrast, is [4][6]: **clear and concise** (plain language, no ambiguity); **complete and accurate** (covers the whole system; following it leads to the expected result); **well-organised and navigable** (a table of contents, descriptive headings, numbered sections, figures, and cross-references — *skimmable*); **audience-aware** (it states who it is for and how to read it); **visual** (diagrams and examples carry what prose cannot); **single-purpose per section** (Diátaxis discipline [1]); **honest** (separates "implemented" from "future work"; states limitations); and **maintainable and up-to-date**. To these standard criteria this document adds one the library itself enforces — **reproducibility**: a reader can regenerate the reported facts by re-running each platform's firewall.

Bad (anti-pattern)	Good (correction)
"The system is fast and secure." (unverifiable adjectives)	"Median API response < 200 ms; content is verified by an automated firewall (§10)." (measurable, cross-referenced)
A tutorial interrupted by reference notes and edge cases.	A tutorial that teaches one path; edge cases live in a separate reference section [1].
"See the diagram" with no figure number.	"As shown in Figure 7.1..." — every figure numbered and referenced in prose.
A metric quoted from memory that no longer matches the code.	A metric measured from the live system at write time and regenerable on demand (§9.4).

Table 2.1 — Concrete anti-patterns and their corrections.

2.3 Maturity: Bare-Minimum to Ultra

Not every project needs the same depth; it is useful to picture a spectrum. This document deliberately targets the **Ultra** column, because it is a printed, examined, hand-to-a-team graduation deliverable.

Dimension	Bad (anti-doc)	Bare-Minimum	Good	Ultra
Existence	None / misleading	A README	Structured report	Multi-view report + living, regenerable artefacts
Accuracy	Fabricated	Mostly right	Verified vs. code	Verified *automatically* — facts reproducible on demand
Structure	Wall of text	Headings	TOC + numbered sections	TOC, figures, tables, cross-refs, traceability matrix
Models	None	One diagram	Key UML views	All four core UML views + data model + architecture
Honesty	Over-claims	Silent on gaps	Lists limitations	Ethics/risk + explicit future work + assumptions
Usability	Unreadable	Readable by author	Readable by team	Audience map; print- & screen-ready; glossary

Table 2.2 — Documentation maturity spectrum.

2.4 The Standard Shape; Standards Adopted

An academic software-engineering report has a conventional shape that readers rely on to orient themselves, and this document follows it: a **title page**, an **abstract** with keywords, a **table of contents**, **lists of figures and tables**, numbered **chapters** proceeding problem → methodology → platforms → analysis → design → implementation → testing → results → conclusion, a **references** section, and **appendices**. This document was produced in conformance with **six recognised documentation and software-engineering standards and frameworks** — being precise, **four are formal international standards** (ISO/IEC/IEEE 29148, IEEE 1016, OMG UML 2.x, and ISO/IEC/IEEE 26514/26515) and **two are established, widely-adopted frameworks/conventions** (the Diátaxis framework and the academic report structure) — together with the project's own content-integrity rule (the "One Law"), which is not an external standard but is stated here for completeness. Table 2.3 names each and states exactly how it was applied and why; the document does not claim conformance to any standard it did not actually follow.

Standard	How it was applied here
ISO/IEC/IEEE 29148 (Requirements Engineering; supersedes IEEE 830 SRS)	The functional and non-functional requirements (§7.4) and the traceability matrix (Appendix A) follow the SRS structure: stakeholders/actors, numbered FR/NFR, and constraints.
IEEE 1016 (Software Design Descriptions)	Chapter 8 presents the design as SDD views — architectural (layered), logical (class & data model), and deployment — each with a diagram and a rationale.
OMG UML 2.x	Behavioural views (Use Case, Activity, Sequence, State Machine) in Chapter 7 and structural views (Class, Deployment) in Chapter 8, drawn in standard notation.
ISO/IEC/IEEE 26514 / 26515 (design and management of user/technical documentation, incl. agile)	The documentation itself was produced with the iterative, continuously-verified process of §1.6, and is structured for its stated audiences (Table 1.2).
Diátaxis framework	Each section holds a single purpose (reference vs. explanation), avoiding the "mixing types" failure (§2.2).
Academic report / thesis structure	Front matter (title, declaration, acknowledgement, abstract, contents, lists) → numbered chapters (problem → analysis → design → implementation → testing → conclusion) → references → appendices, modelled on university final-year-project templates.
The library's own "One Law" (project-specific)	Every quantitative fact reported was *measured* from the live systems and is regenerable, not estimated (Strategy 1, Appendix D).

Table 2.3 — Global documentation standards followed, and how they were applied.

3 The Egyptian Awareness Library

This chapter documents the **mother** of the three platforms: the mission, the governing law, the doctrine, the shared pipeline, and the identity that every platform inherits. Understanding this chapter is a prerequisite for the platform chapters (4–6), because their differences are variations on this common foundation.

3.1 Mission and Vision

Mission. To raise the awareness of Arabic speakers — their ability to verify what is true, to prepare for what is dangerous, and to understand themselves — using content that is verified, cited, culturally situated, and honest about uncertainty.

Vision. A generation of Arabic speakers who can read a claim and ask where it came from; who know what to do in the first minutes of an emergency; and who can tell distress from disorder and a psychological con from real help — all without either panic or false certainty. The library measures its own success not by time-in-app but by whether these competencies rise; it aims, in the long run, to make itself decreasingly necessary.

3.2 The One Law (Anti-Fabrication Doctrine)

The single most important architectural decision across the library is a lesson learned early and applied everywhere: **AI is an *assembler* of grounded content, not a *factory* that generates it.** Formally, **The One Law** states:

> *No factual content claim reaches a user unless it is extracted from a real, fetched source that has been stored and hash-stamped, and every stored quotation appears **verbatim** in that source. A citation that cannot be re-verified is a **critical defect**, not a style issue.*

The law's scope is deliberate: it governs the platforms' **teachable content** — the claims, cards, and answers that assert something about the world — and *not* ordinary interface copy, navigation, or presentation pages, which make no factual assertion to verify.

This inverts the usual risk model of AI content systems. Rather than trusting a model to be truthful and hoping to catch errors, the library **assumes** a model may fabricate and makes fabrication *mechanically impossible to publish*: an assembler may rephrase and sequence sourced material, but it may never introduce a claim, statistic, or citation not already present in a fetched source. The rule is enforced by an automated **firewall** (§3.4, §8.2) that re-reads every source and fails the build if a single quotation is missing.

3.3 The Verification Doctrine

Beneath the One Law sits a philosophy of verification (منهج التثبّت, "the method of ascertaining"). Its operating commitments are: (i) **operational definitions** — a claim must be stated precisely enough to be checked; (ii) **falsification tests** — for any claim, ask what would prove it wrong; (iii) **calibrated verdicts** — separate the empirical, the interpretive, and the value-laden, and never assert 100% certainty where none exists; and (iv) a **harm gate** — content that could cause harm receives stricter scrutiny and routing. These commitments are realised differently per platform but share the same spirit: intellectual humility enforced in software.

3.4 The Shared Content Pipeline

All three platforms turn raw sources into publishable content through the same conceptual pipeline (Figure 3.1). Discovery finds candidate sources; each is fetched and its bytes are hashed (SHA-256) and stored, so the exact evidence is frozen; content is *extracted* (not recalled) as short quotations; each item is *classified* by honesty state and evidence strength; it is *localised* into Arabic and its regional context; and finally the **firewall** re-verifies every quotation against the stored bytes. Anything that fails is rejected and returned for re-grounding; only verified content is published.

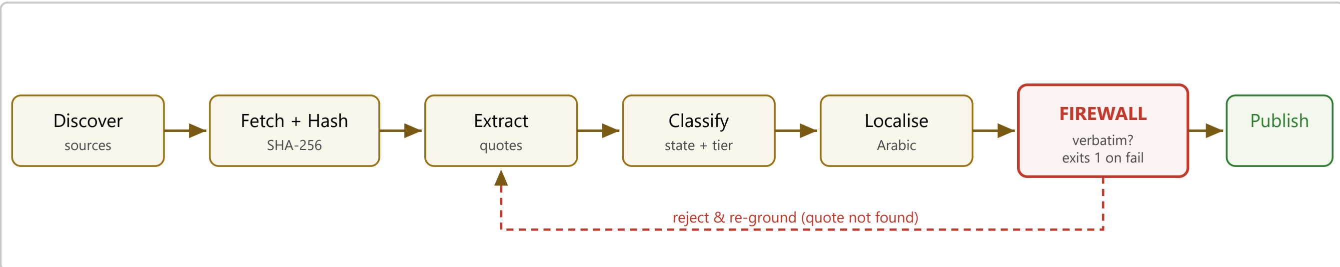

Figure 3.1 — The shared content pipeline. The firewall is a hard gate: it re-reads stored bytes and rejects any item whose quotation is not found verbatim.

3.5 Brand Identity: The Hall of Truth

The library's identity draws on ancient-Egyptian imagery of judgement and truth — the **Hall of Two Truths**, where **Ma'at's feather** weighs the heart, and the **Horus** falcon guards. This is not decoration: it encodes the product promise — that here, claims are *weighed*, not merely asserted. A warm gold palette and the Reem Kufi / Cinzel type pairing give the family a coherent, dignified look, from which each platform then departs to signal its own domain (§8.3).

3.6 The Family of Platforms

The three platforms address the three arenas named in §1.1. They share the pipeline, the firewall, the honesty states, and the Arabic-first, accessibility-conscious design language; they differ in domain, data unit, and visual identity (Figure 3.2, Table 3.1). What distinguishes the whole family from the alternatives a user might otherwise reach is summarised in Table 3.3.

Dimension	Typical wellness app / fact-check site / AI chatbot	The Egyptian Awareness Library
Where content comes from	Author authority, or a model that may generate plausible text.	Extracted from fetched, hash-verified sources; fabrication is unpublishable.
Uncertainty	Hidden or smoothed into confident advice.	Shown as one of four honesty states, with a "what we don't know" card per domain.
Population validity	Ignored; Western findings presented as universal.	A WEIRD flag on every claim marks where it was *not* tested on the user's region.
Language & culture	Machine-translated; culturally generic.	Arabic-first authoring with explicit cultural and religious notes.
Safety	A disclaimer in the footer.	Mechanical crisis routing, non-diagnosis stance, and a clinician gate for high-stakes content.

Table 3.3 — Positioning against typical alternatives.

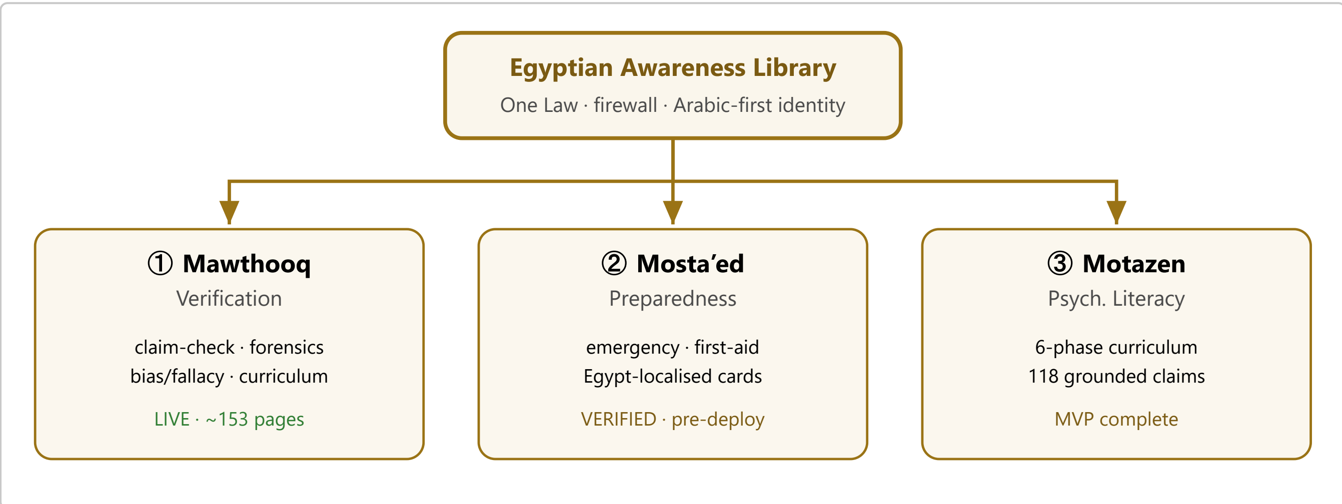

Figure 3.2 — The family of platforms; all three inherit the One Law, the firewall, and the shared identity.

Attribute	Mawthooq (موثوق)	Mosta'ed (مستعد)	Motazen (متزن)
Domain	Verification & critical thinking	Emergency preparedness	Mental-health literacy
Data unit	Claim / engine result	Readiness card	Claim object
Firewall	Engine-level verification + One-Law content	`validate-cards.mjs` (exits 1)	`validate-claims.mjs` (exits 1)
Stack	Next.js 15 / React 19 / TS / Tailwind	Static HTML/JS + Node pipeline	Static HTML/JS + Node pipeline
Identity	Dark "Hall of Truth" gold	Dark preparedness green/compass	Light "breath & paper" sage
Status	Deployed (live)	Validator green; pre-deploy	MVP complete; pre-deploy

Table 3.1 — The three platforms at a glance.

3.7 Source Tiers, Licensing, and the Content Lifecycle

Not all sources are equal, and the library says so. Every source carries a **tier** that is surfaced to the user, so that a claim resting on a single old study is never displayed as if it rested on a consensus of syntheses. The tiers, in descending strength, are shown in Table 3.2.

Tier	Kind of source
A — consensus / synthesis	Systematic reviews and meta-analyses (e.g., Cochrane), large syntheses, and consensus bodies with transparent methodology (WHO, NICE, professional colleges).
B — primary	Well-powered, ideally pre-registered primary studies; post-2015 replication-reform-era work is weighted higher.
C — literacy-specific	The mental-health-literacy and health-communication corpus that grounds *how* the content is taught.
D — critical / negative-science	The replication-crisis literature, methodological critiques, and dissenting reviews — cited deliberately to show both sides of a contested claim.

Table 3.2 — Source tiers.

Licensing. Academic sources are *quoted* — short spans, attributed — under fair-use, and never republished in full; the library stores only the minimal quotation needed to verify a claim, plus the hash and retrieval date. **Lifecycle.** A claim moves through: discover → ingest & hash → extract → classify (state + tier) → gap-fill (find a dissenting source for anything contested, and the domain's "what we don't know") → WEIRD-flag → localise → human review → publish, followed by **periodic re-review**. Because the science moves, every claim carries a `last_reviewed` stamp, and contested claims are re-checked as the field evolves; stale certainty is treated as a failure mode, not a convenience.

4 Platform I — Mawthooq موثوق (Verification)

Mawthooq ("trusted") is the flagship and the most mature platform of the library: a comprehensive verification and critical-thinking web application. Where the other two platforms teach within a bounded domain, Mawthooq is a broad toolkit for interrogating *any* public claim — text, image, or argument — and for training the underlying skill of disciplined thinking.

4.1 Purpose and Users

The purpose is to shrink the population's vulnerability to misinformation by putting verification tools and critical-thinking education directly into users' hands. Its users are ordinary citizens encountering a suspicious claim or image; students and educators using the curriculum; and a small set of administrators and reviewers. Access is **invitation-only** during the current phase, which keeps the reviewed content controlled while the platform matures.

4.2 Feature Set

Mawthooq is large — on the order of **153 rendered pages and 95 API endpoints** (every one is enumerated in Appendix F). Its capabilities group into families (Table 4.1): claim verification, media forensics, cognitive-bias and fallacy detection, a structured critical-thinking curriculum, and supporting infrastructure (accounts, dashboards, certificates). The recurring pattern is the **engine**: a focused analyzer that takes an input (a claim, an image, an argument) and returns a structured, explained assessment — never a bare verdict, always the reasoning and, where relevant, the sources behind it.

Family	Representative modules	What it does
Claim verification	Firewall; Claim-Debunker console; Evidence Search; Decompose	Separates empirical/interpretive/value claims, applies a calibrated verdict, and never asserts 100% certainty.
Media forensics	DeepReal / DeepReal-Forensics; Forensic-C2PA; Forensic-Image; Check-Image	Screens images and media for manipulation and provenance signals (e.g., C2PA content credentials).
Reasoning defense	Bias-Detector; Bias-Fingerprint; Fallacy-Engine; Debate-Sim; Council; Devil-Agent	Names the cognitive bias or logical fallacy at work and rehearses counter-arguments.
Education	Cognition-Curriculum; Critical-Thinking; Epistemology; Mind-Gym; Assessment; Certificate	A multi-day curriculum that trains verification as a skill, with graded assessment and certification.
Domain shields	Arabic-Shield; Drug-Checker; Family-Kit	Targeted protections against region- and topic-specific misinformation.
Infrastructure	Account; Dashboard; Admin; Chatbot; API (~95 routes)	Authentication, progress tracking, administration, and a scoped assistant.

Table 4.1 — Mawthooq: representative feature modules (illustrative of ~153 pages).

A guiding constraint across every engine is the verification doctrine (§3.3): outputs distinguish what is empirically settled from what is interpretive or value-laden, and the platform is designed to build appropriate skepticism rather than to hand down authority. The critical-thinking curriculum includes an **epistemic-firewall** track — advanced, multi-layer exercises in which learners dissect a claim across every level from source to logic to statistics — reflecting the library's conviction that the durable defense against misinformation is a trained mind, not a lookup table.

Beyond the standard engines, Mawthooq ships a set of **flagship capabilities** (Table 4.2) that go well past a typical fact-checker and are, collectively, its strongest competitive differentiators — from a structured eight-layer deception scan to open-source-intelligence investigation, message-forwarding analysis, and demographic-targeted "shields."

Capability	What it does
8-Layer Deception Scan	Dissects a claim across eight layers at once — source, evidence quality, statistics, logic, framing, motive, media integrity, and consensus — the human counterpart to the automated firewall.
Media forensics suite	DeepReal deepfake detection, C2PA content-provenance verification, and reverse-image checks for manipulated media.
OSINT Investigator	Guided open-source-intelligence workflow for tracing a claim, account, or image to its origin.
WhatsApp / message analyzer	Analyses forwarded messages — the dominant misinformation vector in the region — for known rumour and manipulation patterns.
Paper Auditor	Audits an academic paper or a cited study for methodological and citation red flags.
Knowledge Graph & Deception Atlas	Maps claims, entities, and their relationships; live rumour heatmaps and a misinformation atlas track what is circulating now.
Inoculation Passport	A "pre-bunking" track that immunises users against manipulation techniques before they encounter them, recorded as a credential.
Multi-agent debate	Council, swarm-debate, and devil's-advocate agents argue a claim from multiple adversarial angles so weaknesses surface.
No-Hallucination Agent	An assistant architected to refuse fabrication — it answers from verified material or declines.
Targeted shields	Men's, Women's, Family, and Arabic "shields" packaging the defences for demographic- and region-specific misinformation.
Religious-epistemology hub	Handles faith-framed claims with respect and precision (the منهج التثبّت method), neither flattening belief nor letting it override evidence.
Certificate & leaderboard	Gamified, graded critical-thinking certification with cross-user competition — motivation without dark patterns.

Table 4.2 — Mawthooq flagship capabilities (the strongest, differentiating features).

4.3 Architecture and Access Control

Mawthooq is built on **Next.js 15** with **React 19** and TypeScript, styled with Tailwind, and deployed on Vercel. It uses the Next.js App Router, co-locating server-rendered pages and API route handlers. Access is gated by an **invitation-code** mechanism: a set of issued codes admits reviewed users, keeping the audience controlled during the current review phase. A themeable design layer supports many visual modes bound to design tokens, so that every page honours the accessibility requirement that meaning never depends on colour alone.

4.4 Status

Mawthooq is **deployed and live**. Its verification engines and content pages are in active use, and its automated checks pass on the production build. As the most complete member of the family, it establishes the patterns — the engine shape, the calibrated verdict, the token-bound theming, the invitation gate — that the other platforms adapt to their own domains.

4.5 The Engine Pattern and the Calibrated Verdict

The unit of composition in Mawthooq is the **engine**: a self-contained analyzer with a narrow responsibility. An engine accepts an input, performs one kind of analysis, and returns a *structured, explained* result. This pattern is what lets a platform of roughly 153 pages remain comprehensible — each page is a thin surface over one engine, and engines can be combined without entangling their logic.

What every engine shares is the **calibrated verdict**, the practical form of the verification doctrine (§3.3). Rather than collapsing a messy claim into "true" or "false," an engine first *separates* the claim into its empirical, interpretive, and value-laden parts; assesses each on its own evidence; and returns a verdict that is explicit about its own confidence and about what would change it. The platform is deliberately built so that it **never asserts 100% certainty**, because certainty is exactly the tell of the misinformation it exists to counter. A user leaves an engine not merely with an answer but with a *model of how the answer was reached* — which is the transferable skill.

A second shared property is a **harm gate**: engines that touch sensitive territory (health, self-harm, legally or medically consequential claims) apply stricter handling and route to appropriate resources rather than adjudicating alone. As with the other platforms, Mawthooq is a tool for thinking, not a substitute for a professional.

4.6 The Critical-Thinking Curriculum

Beyond point tools, Mawthooq treats verification as a **skill to be trained**, not merely a service to be consumed — reflecting the library's conviction that the durable defence against misinformation is an educated mind. Its curriculum runs over many structured days and includes an advanced **epistemic-firewall** track in which learners take a single claim apart across every layer at once — source, evidence quality, statistics, logic, framing, and incentive — the same layers the automated firewall enforces on content, now taught as a human capability. Progress is graded, and completion is recognised with a certificate, giving the abstract goal of "thinking better" a concrete, motivating structure. This curriculum is the clearest expression of the library's long-term ambition: to make itself decreasingly necessary by leaving behind users who no longer need it.

5 Platform II — Mosta'ed مستعد (Preparedness)

Mosta'ed ("prepared") applies the library's method to a different arena: what to do in the first minutes of an emergency. Its content unit is the **readiness card** — a compact, source-verified guide to a specific situation (a burn, a choking, a panic attack, a suspected heart attack), written for a frightened non-expert and localised precisely to Egypt.

5.1 Purpose and Users

The purpose is to close a preparedness gap: people often do the wrong thing in an emergency not from indifference but from confidently-held folk beliefs (for example, applying the wrong first-aid to a burn or a poisoning). Mosta'ed replaces folk belief with verified, actionable steps drawn from recognised public-health sources. Its users are ordinary members of the public; the content is authored and reviewed under strict provenance controls before it can be shown.

5.2 The Card Provenance Pipeline

Each card is a structured object whose actionable steps — the "do," the "don't," the red-flags, and whom to call — are not written from memory but **extracted** from fetched, hash-stamped source documents. Every step carries a source sentence and its Arabic rendering. A dedicated validator, `validate-cards.mjs`, enforces the One Law for this platform: it **exits with a non-zero code** on any hard failure, and its load-bearing rule is that a card of high severity (level 4–5, i.e., potentially life-critical) with *zero* ingested sources is a hard failure — "model recall wearing a citation." A plausible-looking URL is explicitly *not* grounding; a source counts only if it was actually fetched and hashed into the corpus. In its current state the platform reports **all cards grounded** against a corpus of hundreds of hash-verified documents, with the validator passing (exit 0). Content is produced in gated "waves," each batch adversarially reviewed before it may ship, and several early fabrications (an incorrect poisoning instruction, a wrong medication dose) were caught and corrected by exactly this gate — a concrete demonstration of the One Law preventing harm.

Each readiness card is a structured object whose parts map to what a frightened person actually needs, in order. Table 5.1 shows its anatomy; every actionable line carries a source sentence and its Arabic rendering, so the whole card is verifiable field by field.

Part	Purpose
Title & severity	The situation and its level (L1–L5); level 4–5 (life-critical) triggers the strictest grounding and clinician-review rules.
When / why	How to recognise the situation this card is for.
Do (steps)	The correct actions, in order — each grounded in a fetched source sentence.
Don't	The dangerous folk-remedies to avoid (often the whole point of the card).
Red flags	Signs that mean "escalate now."
Who to call	The correct *Egyptian* number (123 / 122 / 180 / 16328), never a foreign one (Gate 6).

Table 5.1 — Anatomy of a readiness card.

5.3 Egyptian Localisation (Gate 6)

A safety rule unique to this platform, "Gate 6," governs geography: **foreign emergency numbers must never reach an Egyptian reader.** Source documents frequently instruct readers to "call 911" or "dial 999"; the pipeline drops these and substitutes Egypt's own numbers in the reader-facing text — **123** (ambulance), **122** (police), **180** (civil defence / fire), **112** (unified emergency), and **16328** (mental-health support). Foreign numbers survive only inside code comments and inside the retained source-quotation fields required for verbatim verification; they never appear in the Arabic instructions a user reads. Where no verified national line exists for a hazard, the card says so honestly and routes to the ambulance.

5.4 Status

Mosta'ed's content engine is built and **honest**: the validator is green and localisation is clean. Its remaining work is breadth (more domains beyond the emergency/first-aid core) and two standing launch blockers shared with Motazen: a licensed-clinician sign-off on the highest-severity cards, and deployment. It is documented here as a **verified, pre-deployment** platform.

6 Platform III — Motazen متزن (Psychological Literacy)

Motazen ("balanced") turns the library's method inward, onto the human interior. It teaches Arabic-speaking users *how to think about* mental health — never what to feel, and never what to do about a specific diagnosis. Its tagline sets the tone: *"Your journey to inner balance begins with a step and a deep breath."* It is the platform built most recently and is documented here in the most detail, as its curriculum is the clearest illustration of the library's pedagogy.

6.1 Purpose and Users

The purpose rests on the academic construct of **Mental Health Literacy** (MHL): the knowledge and beliefs that aid recognition, management, and prevention of mental states, together with the ability to evaluate claims about the mind. The problem it addresses is that the mental-health treatment gap is, first, a literacy gap, worsened by stigma and by a market of confident falsehoods. Its users are ordinary Arabic-speaking adults; it is explicitly *not* for diagnosis or treatment, and it foregrounds routing to human help.

6.2 The Six-Phase Curriculum

Motazen's content library holds **118 grounded claims** supported by **129 sources** (130 fetched, hash-stamped documents), sequenced as a "descending-then-ascending" journey from a conceptual foundation to a defined end state:

- **Phase −1 — "Below Zero" (40 claims).** The primer nobody teaches: what mental health even *is*; whether "mental" is separate from "physical"; why the topic needs science; how we know something is true in psychology (replication, WEIRD); normal vs. distress vs. disorder; who decides what counts as a disorder. It ends with a four-part exit criterion the user should be able to state in their own words.
- **Phase 0 — "Map of the Territory" (6 claims).** The conceptual map: the biopsychosocial model taught as a useful *model* (with its critique), and the two-axes insight that mental health is not merely the absence of illness.
- **Phase 1 — The Domains (63 claims).** Six highest-burden domains, each following the same shape — recognition, established, contested, "what we don't know," normal-vs-disorder, the regional caveat, and the misinformation cluster: **mood/depression (14), anxiety (12), stress & sleep (12), trauma (9), attention/ADHD (8),** and **the psychotic spectrum (8)**.
- **Phase 2 — Skills (5 claims).** Evidence-tiered self-help skills (slow breathing, behavioural activation, gradual exposure, gratitude), each with its "this is self-help literacy, not treatment" boundary.
- **Phase 3 — Navigation (4 claims).** Help-seeking in Egypt, a five-question claim-evaluation lens, and family/cultural-religious guidance.
- **Phase 4 — "Balance."** A self-assessment of the literate end-state competencies (see §6.4).

6.3 The Nine Surfaces

The library is delivered through nine interactive surfaces, each a *view* onto verified claim objects — none generates content. Following the product's scope discipline (cards, games, scenarios, exercises, and a chatbot — "nothing more"), they are: the **descent journey** (Phase −1), the **map** (Phase 0), the **domains** (Phase 1), the **exercises** surface (Phase 2), the **navigation** surface (Phase 3), the **self-assessment** (Phase 4), a status-judgment **game** ("Judge for Yourself"), everyday **scenarios**, and a scoped **assistant**. The assistant is deliberately the most constrained surface: it is **retrieval-only** (it matches a query against the library and cannot free-generate psychological claims), it refuses diagnosis and medication questions, and on any signal of crisis it drops the literacy role and routes to human help — a non-bypassable gate (§7.3).

6.4 Safety and the Four Honesty States

Every claim carries exactly one **honesty state — established, contested, debunked,** or **unknown** — which the interface is structurally forbidden to hide. Across the library these currently number **75 established, 19 contested, 16 unknown,** and **8 debunked**. A **contested** claim must store a real dissenting source and is shown with both sides; every domain must carry a "what we don't know" card. Each claim also carries a **WEIRD flag** indicating whether the finding was tested on populations resembling the user, and a **crisis-sensitive** flag (set on 20 claims) that triggers safety handling. A persistent crisis banner and the Egyptian support numbers (16328, 123) appear across the surfaces, and the judgment game excludes crisis-sensitive content by design. These are not stylistic choices; they are enforced requirements (§7, §10).

6.5 The Six Domains in Detail

Phase 1 is the largest part of the curriculum. Each of its six domains is authored to the same seven-part shape — *recognition* (what it feels like from the inside), *established* facts, *contested* questions (shown with a dissenting source), a *"what we don't know"* card, the *normal-versus-disorder* line, the *regional/WEIRD caveat*, and the *misinformation cluster* to inoculate against. Table 6.1 inventories the domains; the discipline of the shape is what keeps a high-stakes topic honest rather than either alarmist or falsely reassuring.

Domain	Claims	Representative teaching (with an honesty state)
Mood / Depression	14	Depression is more than sadness (*established*); the "chemical-imbalance/serotonin" story is *debunked*; how clinically meaningful antidepressants are is *contested*; long-term psychotherapy durability is *unknown*; "depression is weakness" is *debunked*.
Anxiety & Fear	12	Fear is an adaptive alarm (*established*); CBT/exposure is effective (*established*); publication bias inflates the apparent drug evidence (*contested*); why some develop a disorder is *unknown*; the CBD/cannabis "cure" is *debunked*.
Stress & Sleep	12	Chronic stress causes cumulative wear (*established*); CBT-I is first-line for insomnia (*established*); "8 hours for everyone" is *contested*; why we sleep at all is *unknown*; "cortisol-detox" supplements are *debunked*.
Trauma & PTSD	9	Most people exposed to trauma do *not* develop PTSD (*established*); trauma-focused CBT has the best evidence (*established*); single-session "debriefing" is *debunked* (may harm); EMDR's superiority is *contested*; the displaced are under-studied (*unknown*).
Attention / ADHD	8	ADHD is a real, valid condition (*established*); over-/under-diagnosis is *contested*; exact causes are *unknown*; "sugar / bad parenting / screens cause ADHD" is *debunked*.
Psychotic Spectrum	8	Psychotic experiences lie on a continuum (*established*); the "violent and dangerous" stereotype is *debunked* (they are more often victims); how long to maintain medication after a first episode is *contested*; long-term functional outcomes are *unknown*. Handled with maximum dignity.

Table 6.1 — Phase-1 domain inventory (63 claims).

The recurring presence of a *debunked* misinformation card and an *unknown* card in every domain is deliberate: it is the platform's "negative science" in action — teaching not only what is known but how knowledge fails, and training users to expect and tolerate genuine uncertainty rather than flee from it into false certainty. This is the graduate-level skill the platform exists to build.

6.6 Example Claims (as Published)

To make the abstract concrete, four verified claims are reproduced below exactly as the library stores and displays them — each with its honesty state, its verbatim source quotation, and its English gloss. These are not illustrations invented for this document; they are live records that pass the firewall.

✕ **Debunked** strong evidence against

«الاكتئاب مجرّد نقص سيروتونين، والحبّة بتظبّطه» — القصّة البسيطة دي مالهاش دليل متّسق.

English: The simple story — "depression is just a serotonin deficiency the pill fixes" — has no consistent evidence.

Source (verbatim): "no support for the hypothesis that depression is caused by lowered serotonin activity or concentrations" — Moncrieff et al., *Molecular Psychiatry*, 2022.

✓ **Established** strong

الصحة النفسية مش عكس المرض — دول محورين مختلفين، مش طرفين لخطّ واحد.

English: Mental health is not the opposite of mental illness — they are two distinct dimensions, not two ends of one line.

Source (verbatim): "a distinction between the absence of mental illness and positive mental well-being" — VanderWeele et al., 2026.

╳ **Debunked** anti-stigma

«اللي عنده ذُهان خطر وعنيف» — غلط؛ هم للإيذاء أقرب، ضحايا أكتر ما هم جُناة.

English: "People with psychosis are dangerous and violent" is false; they are more often victims than perpetrators.

Source (verbatim): "People with severe mental illness (SMI) face a significantly higher risk of victimization than the general population" — Krumm et al., 2026.

⇄ **Contested / self-help** mixed evidence

التنفّس البطيء ممكن يساعد على التوتّر — أداة مفيدة، مش معجزة.

English: Slow breathing may help with stress — a useful tool, not a miracle. (Carries a self-help-not-treatment boundary.)

Source (verbatim): "breathwork may be effective for improving stress and mental health" — Fincham et al., *Scientific Reports*, 2023.

Each box corresponds to one row of the claim table (Table 8.1) and one node of the class diagram (Figure 8.2): a Claim, its honesty state, and a Quotation pointing at a hash-verified Source. The uniformity is the point — the same structure carries a reassuring fact and an uncomfortable one, so the user is never guessing which is which.

7 System Analysis and Requirements

This chapter analyses the system: it models behaviour using the standard UML analysis views and states the system's functional and non-functional requirements. Because the three platforms share one pipeline and one set of roles, the models are drawn at the **library level**, with platform-specific notes where behaviour differs. The four core UML views required — Use Case, Activity, Sequence, and Class — appear here (Use Case, Activity, Sequence, State) and in Chapter 8 (Class); the requirements are stated in §7.5 and traced to design and test in Appendix A.

7.1 Use Case Diagram and Specifications

Figure 7.1 shows the actors and the principal use cases. The **Citizen/Learner** is the primary human actor; **Content Author**, **Clinician/Reviewer**, and **Administrator** are supporting human actors; and the **Firewall** is a secondary *system* actor that participates in every content-publishing use case. Two relationships are highlighted: authoring content **«includes»** validating it (nothing is authored without passing the firewall), and asking the assistant may **«extend»** into crisis help when a risk signal is detected.

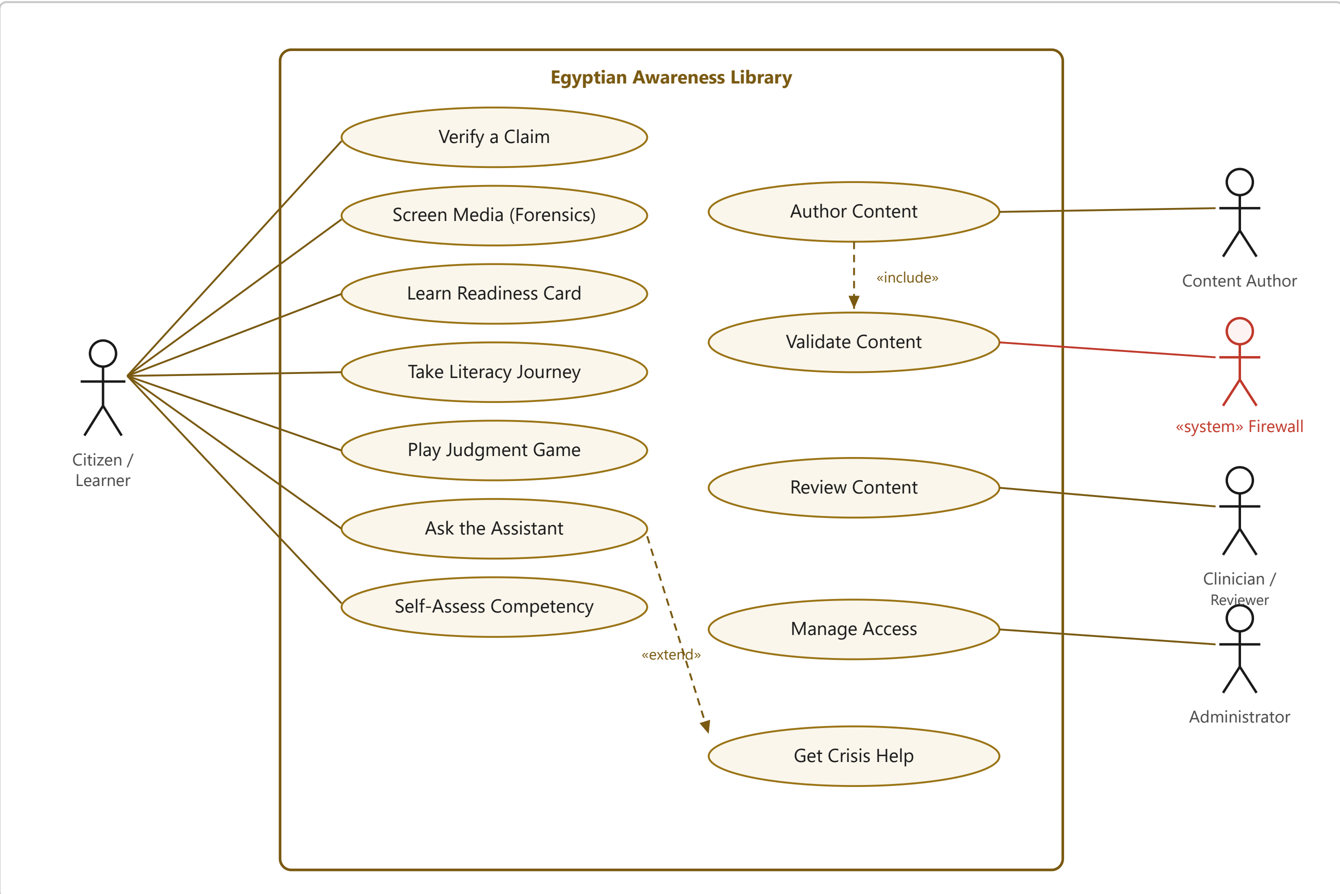

Figure 7.1 — Use Case Diagram for the Egyptian Awareness Library.

Actor	Responsibility
Citizen / Learner	The primary user: verifies claims, learns readiness cards, follows the literacy journey, plays the game, asks the assistant, and self-assesses.
Content Author	Assembles content from fetched sources; cannot publish without passing the firewall.
Clinician / Reviewer	Reviews high-stakes content (crisis-sensitive, high-severity) before public release.
Administrator	Manages access (e.g., invitation codes) and monitors the platform.
«system» Firewall	A non-human actor that re-verifies every quotation against stored bytes; participates in every publish path.

Table 7.1 — Actors and responsibilities.

Field	Description
Actor	Citizen / Learner
Precondition	The claim library is loaded.
Main flow	1. User submits a question. 2. The crisis filter screens the text. 3. If safe, the retrieval engine matches the question against the claim library. 4. Grounded claim cards (with status, tier, source) are returned, or an honest "I don't know."
Alternate flow «extend»	2a. If a self-harm/crisis signal is detected, the assistant abandons retrieval, displays crisis routing (16328 / 123 / a human), and stops.
Postcondition	The user receives only source-verified content or a safe crisis referral; no free-generated psychological claim is ever emitted.

Table 7.2 — Use case specification: "Ask the Assistant" (Motazen).

Field	Description
Actor	Citizen
Precondition	The user has a claim (text) or media they doubt.
Main flow	1. User submits the claim. 2. The relevant engine decomposes it into checkable parts. 3. Each part is assessed and separated into empirical / interpretive / value components. 4. A *calibrated* verdict is returned with its reasoning and, where available, the sources — never a bare "true/false," never 100% certainty.
Alternate flow	2a. For an image, the media-forensics engine screens for manipulation and provenance (e.g., C2PA) signals instead.
Postcondition	The user has a reasoned, sourced assessment and can see *how* it was reached.

Table 7.3 — Use case specification: "Verify a Claim" (Mawthooq).

Field	Description
Actor	Content Author (primary); Firewall, Clinician (secondary)
Precondition	Candidate sources have been discovered.
Main flow	1. The author fetches and hash-stamps each source. 2. Short quotations are extracted (not recalled). 3. Each claim is classified (state + tier) and localised. 4. **«include»** the firewall re-verifies every quotation. 5. On pass, the content enters the library.
Exception flow	4a. If any quotation is not found verbatim, the build fails (exit 1); the claim is rejected and returned for re-grounding. Nothing is published.
Business rule	High-severity / crisis-sensitive content additionally requires clinician review before public release.
Postcondition	Only provenance-verified content reaches users; fabrications are structurally excluded.

Table 7.4 — Use case specification: "Author & Publish Content" (all platforms).

7.2 Activity Diagrams

Figure 7.2 models the **One-Law content pipeline** as a UML activity: the flow that every claim or card follows from a raw source to publication. Its decisive element is the guard on the firewall — *is the quotation found verbatim in the stored bytes?* — whose "no" branch returns the item for re-grounding rather than letting it through. This single guard is what makes fabrication unpublishable.

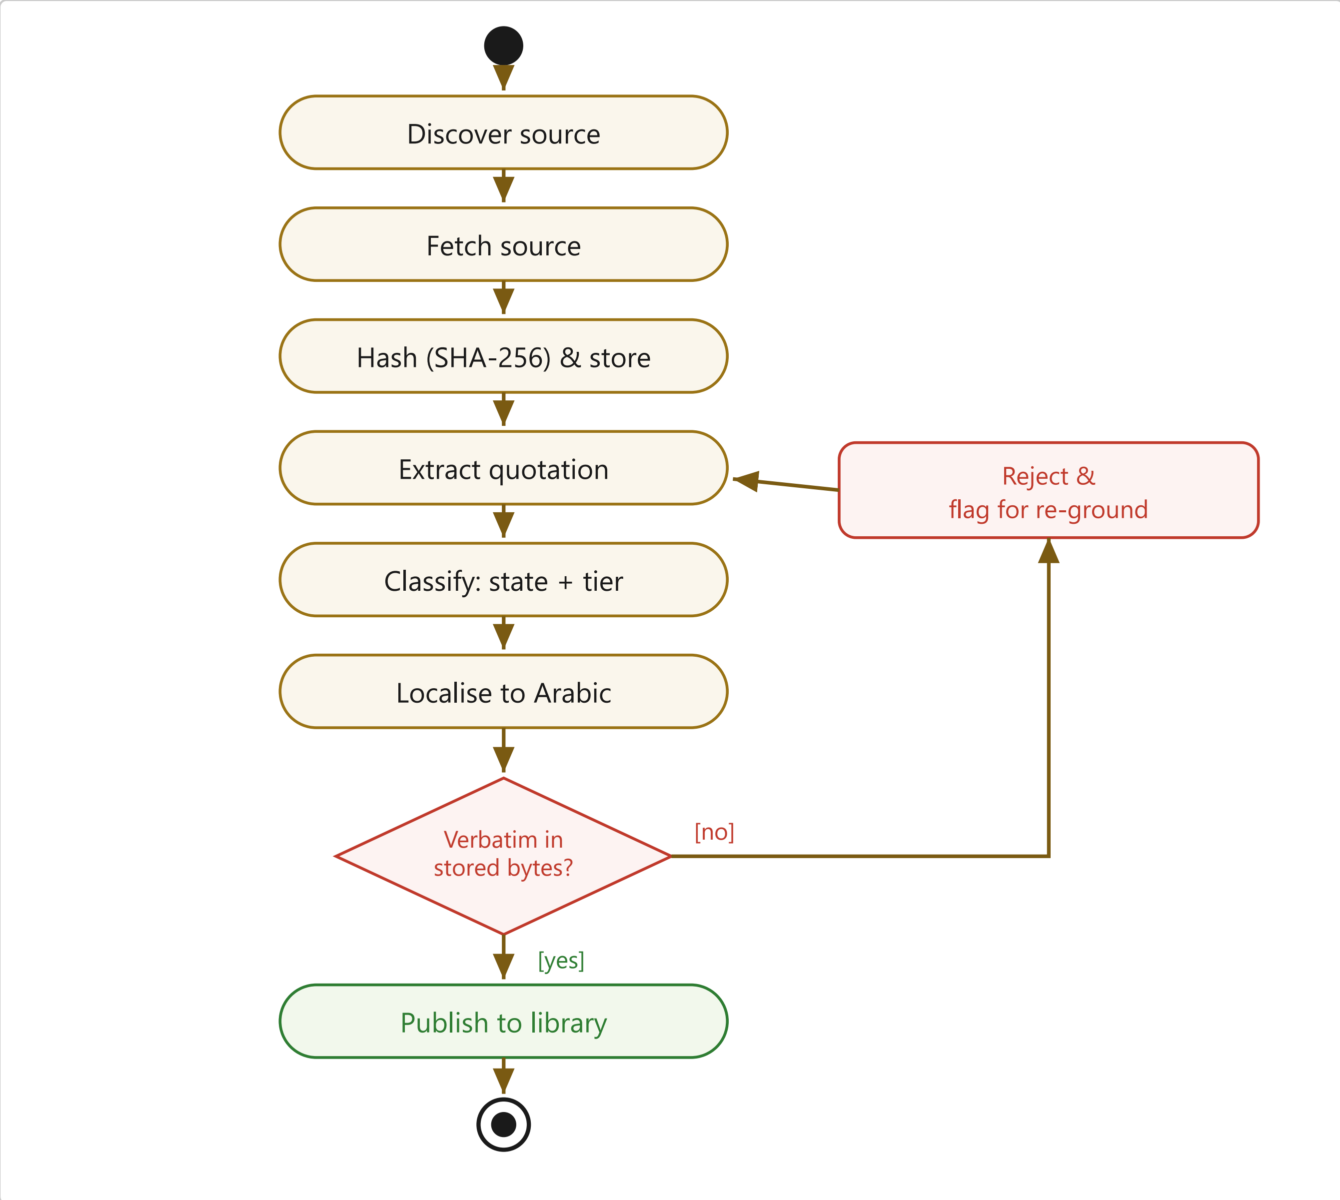

Figure 7.2 — Activity Diagram: the One-Law content pipeline. The firewall guard's "no" branch prevents fabrication from being published.

Figure 7.3 models the **learner's journey** through Motazen's six-phase curriculum: a directed progression from the "below-zero" primer to the self-assessment, with a guard that returns a learner who has not yet met the end-state competencies to revisit the domains.

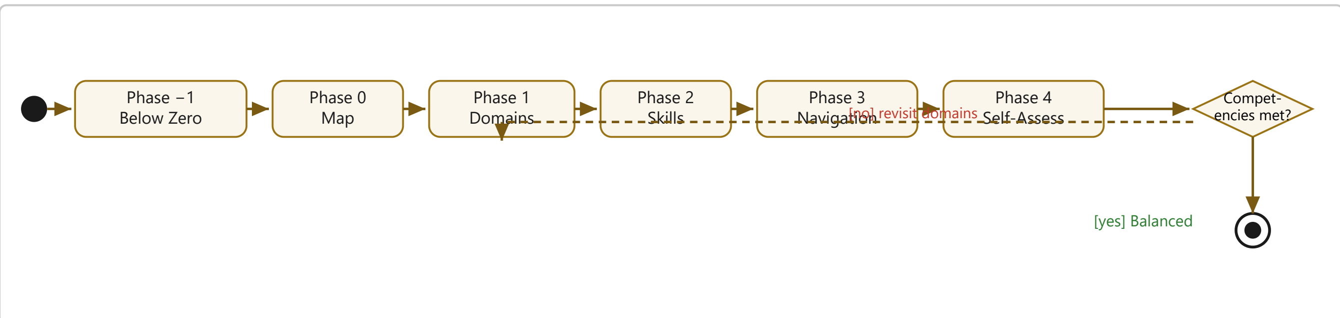

Figure 7.3 — Activity Diagram: the learner's journey through Motazen's curriculum.

7.3 Sequence Diagrams

Figure 7.4 traces the safety-critical interaction: a user asking the assistant a question. The **alt** fragment shows the non-bypassable branch — if the crisis filter detects a risk signal, the assistant abandons retrieval and routes to human help; only in the safe branch does it consult the library and return grounded cards. No path allows the assistant to invent a psychological claim.

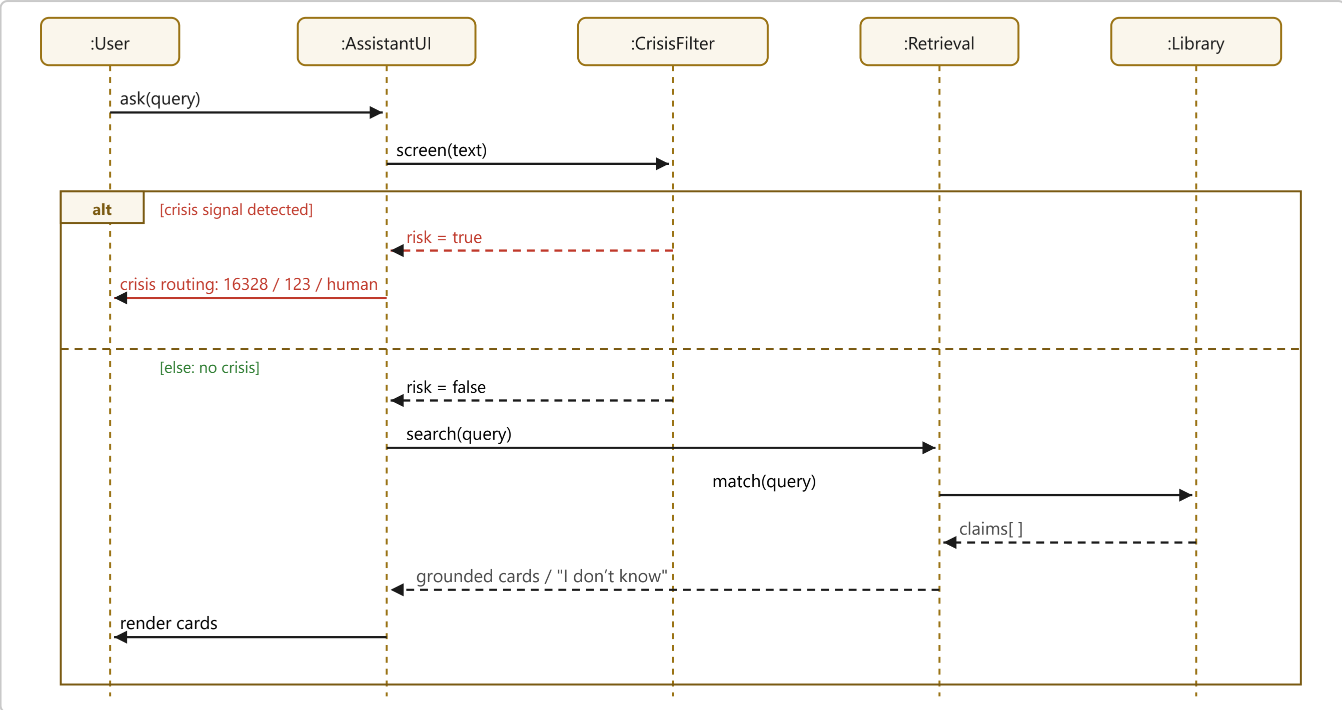

Figure 7.4 — Sequence Diagram: assistant query with non-bypassable crisis routing.

Figure 7.5 traces content **verification** — the firewall in action. The assembler fetches and hashes a source, extracts a quotation, and submits the claim; the validator re-reads the stored bytes and tests whether the quotation is present verbatim. The **alt** fragment is the whole point: a match publishes; a miss fails the build (exit 1) and the claim is dropped.

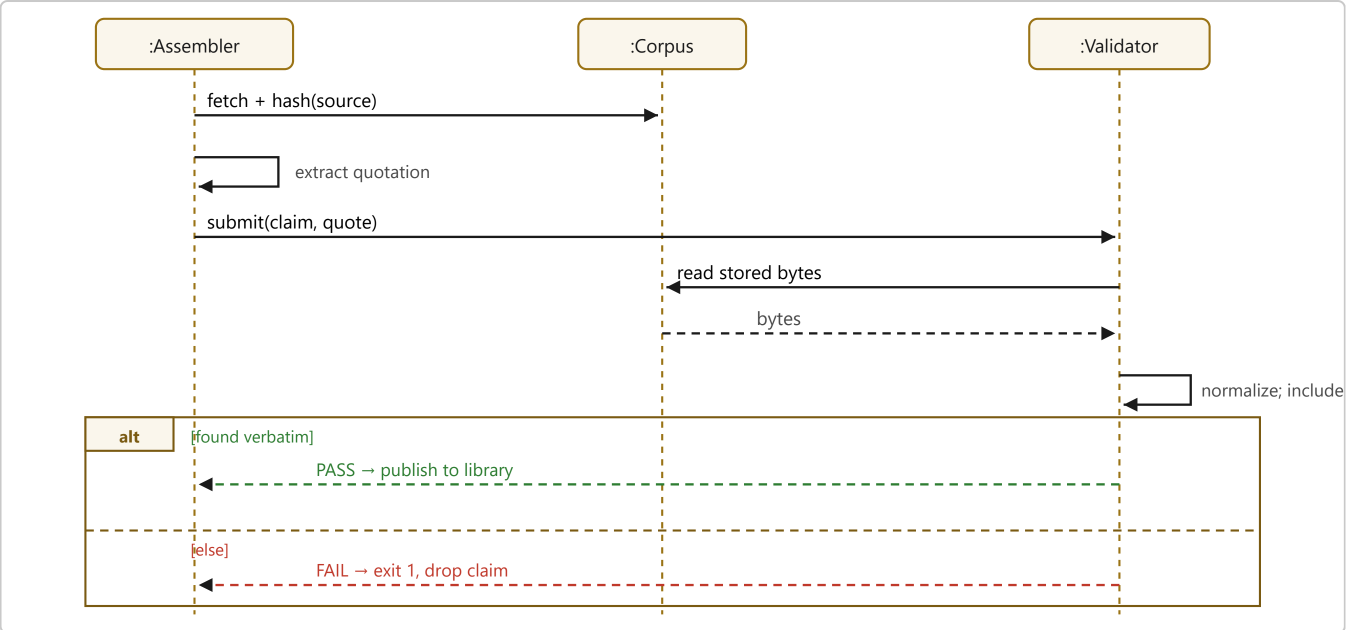

Figure 7.5 — Sequence Diagram: content verification by the firewall.

7.4 Claim Lifecycle (State Machine)

Completing the behavioural view, a single claim moves through a well-defined lifecycle, shown as a UML state machine in Figure 7.6. The pivotal transition is out of *Verifying*: a claim is **Published** only on a verbatim match, and is otherwise **Rejected** back to *Draft* for re-grounding. Because sources go stale, a published claim can re-enter verification through *Under Re-review* as its recency stamp lapses — so "published" is never "finished."

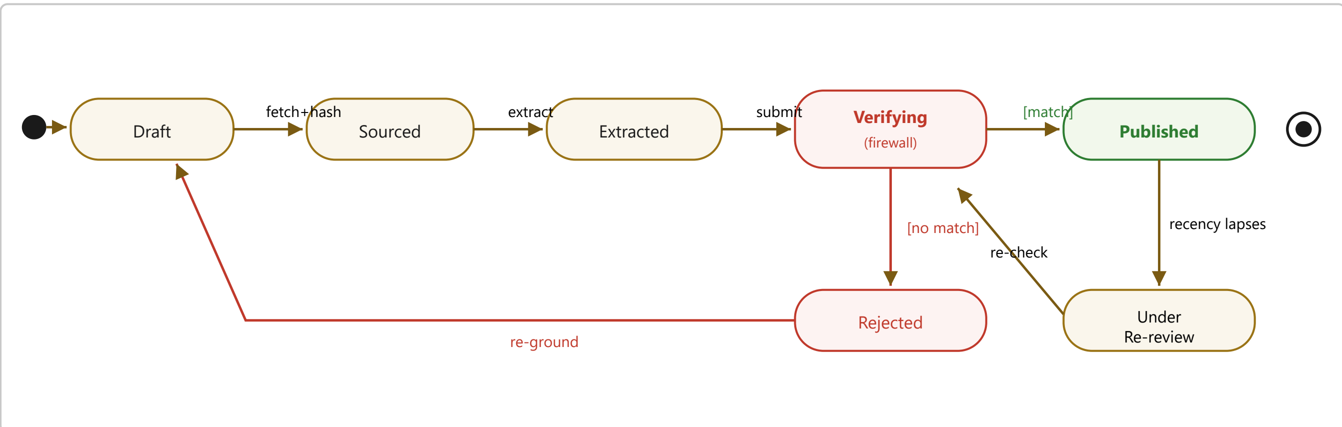

Figure 7.6 — State machine of a claim: only a verbatim match reaches Published; staleness returns it to verification.

7.5 Functional and Non-Functional Requirements

The behavioural models above are realised by the following requirements. **Functional requirements** (FR) state what the system must do; **non-functional requirements** (NFR) state the quality attributes it must exhibit. Each is traced to its design and test in Appendix A.

ID	The system shall…
FR1	display source-verified content with its honesty state, evidence tier, and WEIRD flag.
FR2	enforce the One Law — publish nothing whose quotation is not found verbatim in a hash-verified source.
FR3	classify every claim as established, contested, debunked, or unknown.
FR4	provide verification tools (claim-checking, media forensics, bias/fallacy detection) — Mawthooq.
FR5	provide Egypt-localised, source-verified preparedness cards — Mosta'ed.
FR6	deliver the six-phase literacy curriculum through nine surfaces — Motazen.
FR7	answer user questions with a retrieval-only assistant that cannot free-generate claims.
FR8	detect crisis signals and route to human help on a non-bypassable path.
FR9	let a user self-assess their literacy competencies.
FR10	control access via invitation codes during the review phase (Mawthooq).

Table 7.5 — Functional requirements.

ID	Quality attribute
NFR1	**Integrity:** every quotation must be re-verifiable; fabrication is a build-failing defect.
NFR2	**Usability (Arabic-first, RTL):** the interface is right-to-left and readable on inexpensive phones.
NFR3	**Accessibility:** meaning is never colour-alone; WCAG-AA contrast; reduced-motion respected.
NFR4	**Safety:** non-diagnosis stance, crisis routing, no harmful methods, medication caveats.
NFR5	**Performance:** content-first platforms are static and cache-busted for fast, cheap delivery.
NFR6	**Portability & theming:** self-contained files; legible in light and dark; print-ready.
NFR7	**Maintainability:** idempotent pipeline; recovery in one command.
NFR8	**Privacy:** the content-first platforms (Mosta'ed, Motazen) collect no personal data and keep local state only; Mawthooq collects only the minimal account data required for authentication, progress, and access control.
NFR9	**Reproducibility:** every reported fact is regenerable by re-running the firewall.

Table 7.6 — Non-functional requirements.

8 System Design

8.1 Layered Architecture

Every platform in the library follows the same layered architecture (Figure 8.1). At the bottom, an **acquisition pipeline** fetches and hash-stamps sources into a frozen **corpus**. Above it, the **verification layer** (the firewall) is a hard gate through which content must pass. The verified content forms the **library** — the single source of truth. The **presentation layer** (the surfaces) only ever *reads* the library; it can display nothing the library cannot cite. This strict, one-directional dependency — surfaces depend on the library, the library depends on the firewall, the firewall depends on the corpus — is what guarantees the One Law at the level of architecture, not merely convention.

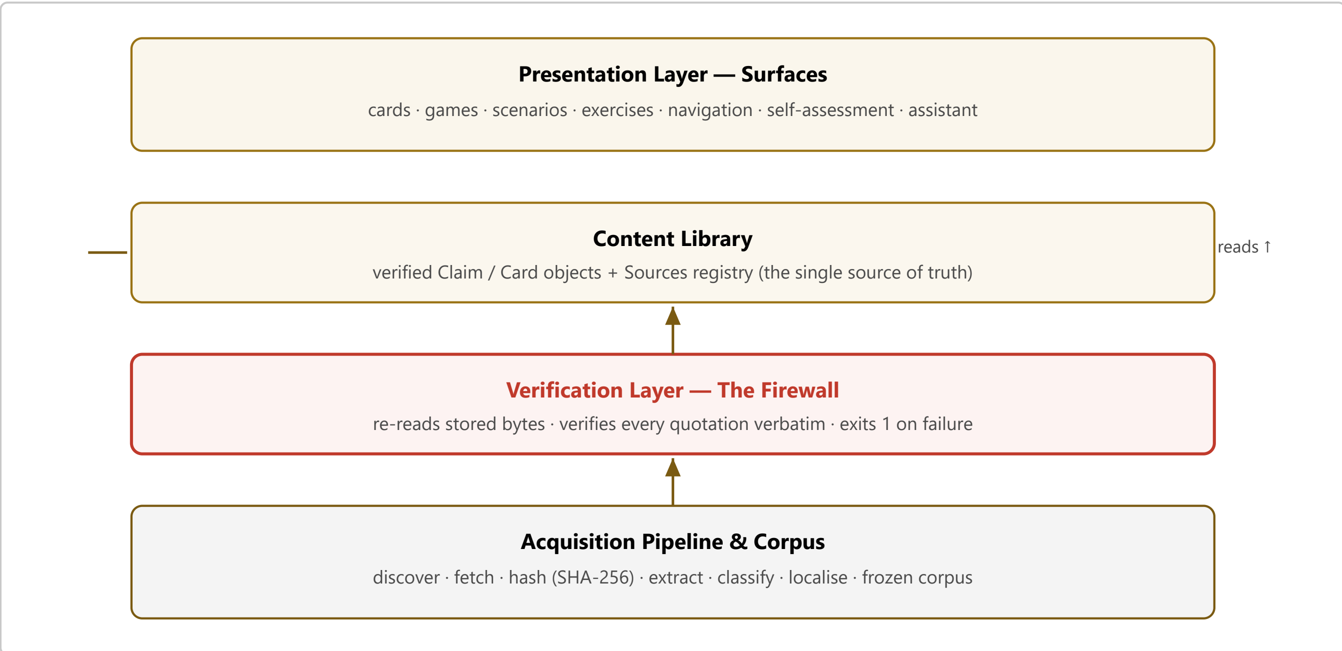

Figure 8.1 — Layered architecture; dependencies point downward, so surfaces can never bypass the firewall.

8.2 Class Diagram and Data Model

Figure 8.2 gives the shared content model as a UML class diagram. The heart of the system is the **Claim** (called a **Card** in Mosta'ed): a bilingual, classified teaching unit. A Claim is composed of one or more **Quotations**, each of which references exactly one **Source** — this composition *is* the provenance chain. The **Firewall** depends on both Claim and Source (it re-reads the Source's stored bytes to verify the Claim's Quotations). A **Platform** aggregates many Claims and composes many **Surfaces**, each of which renders Claims.

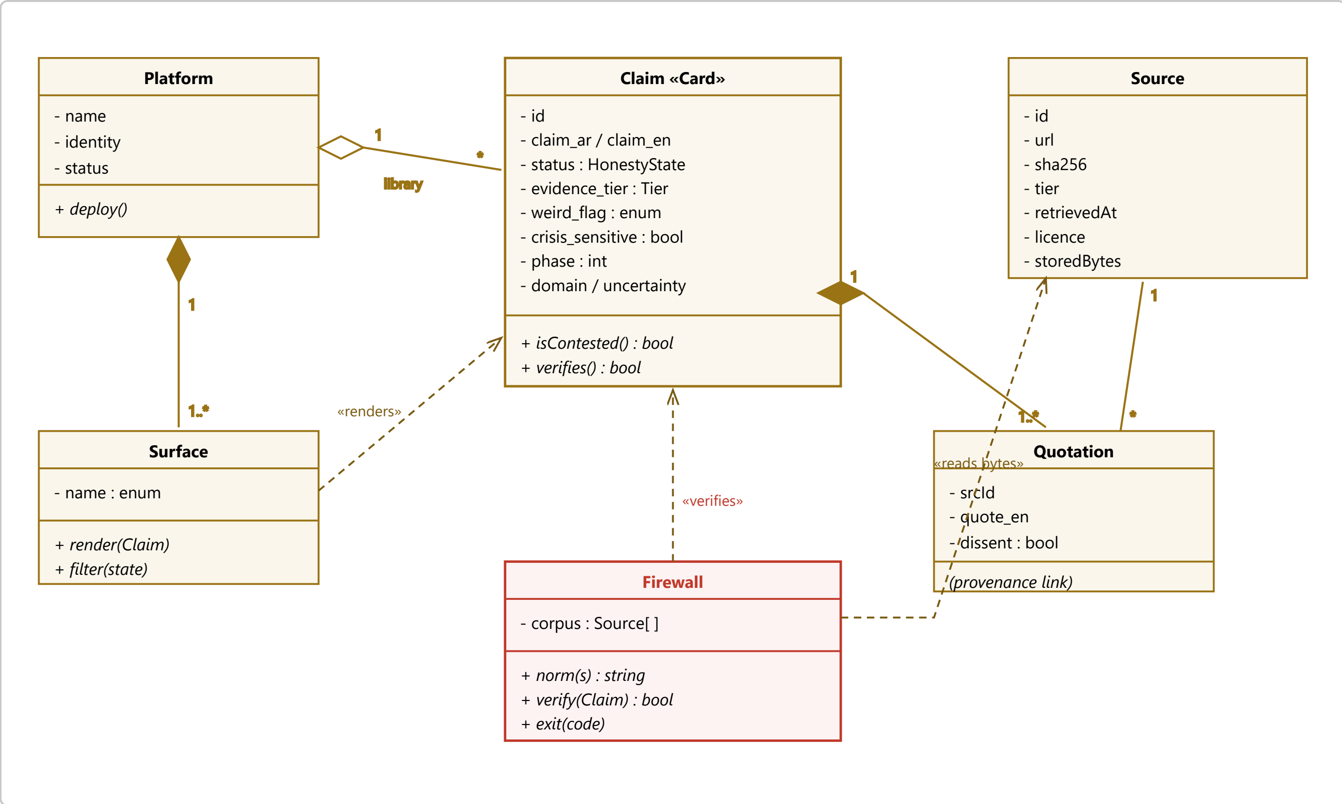

Figure 8.2 — Class Diagram: the shared content model. The Claim–Quotation–Source composition is the provenance chain the Firewall verifies.

Field	Type	Purpose
`claim_ar / claim_en`	string	The statement, bilingual (Arabic primary, English gloss).
`status`	enum	Honesty state: *established* · *contested* · *debunked* · *unknown*. Never hidden by the UI.
`evidence_tier`	enum	Source strength: strong · mixed · weak · debunked · open.
`sources[ ]`	Quotation[ ]	≥1 provenance link; a *contested* claim must include a dissenting source.
`weird_flag`	enum	yes / no / unknown — was this tested on populations like the user?
`crisis_sensitive`	bool	Triggers safety handling; excluded from the game surface.
`cultural_note_ar`	string?	Localisation / religious-epistemology note where relevant.
`uncertainty_note_ar`	string	What we do *not* know about this claim.
`phase · domain`	int · string	Curriculum placement.

Table 8.1 — The Claim/Card object: field dictionary.

8.3 Visual Identity and UX

All three platforms are **Arabic-first** and right-to-left (RTL), and all inherit the same design commitments while expressing distinct identities so a user always knows which platform they are in. **Mawthooq** wears the dark "Hall of Truth" gold, invoking judgement and weighing. **Mosta'ed** uses a dark preparedness palette (a compass/green motif) signalling readiness and calm under pressure. **Motazen** deliberately departs from both with a light "breath & paper" identity — a warm, book-like paper ground and a single deep-sage accent — because a mental-health surface should down-regulate arousal, not shout.

Typography pairs a humanist Arabic face (IBM Plex Sans Arabic, with a serif for long-form and citations) with a matching Latin face, chosen for readability at small sizes on inexpensive phones. **Accessibility** is treated as a requirement: meaning is never carried by colour alone — every honesty state pairs a colour with an icon and a label — contrast targets WCAG AA, motion is breath-paced and respects `prefers-reduced-motion`, and every page is theme-bound so it is legible in both light and dark. Because content is shareable, cards are designed to travel through social feeds as a corrective to the bad cards already circulating there.

8.4 Deployment and Maintenance View

The two architectures imply two deployment stories. **Mawthooq**, as a full application, is deployed on a managed platform (Vercel) with builds triggered on change; it is currently live. The **content-first platforms** compile to **self-contained static files** — the HTML surfaces, a single data file, and the hashed corpus — which can be served from any static host or CDN with no server and no database. This makes them cheap, fast, trivially cacheable, and easy to audit: the entire published system is a handful of readable files. Figure 8.3 shows the deployment topology.

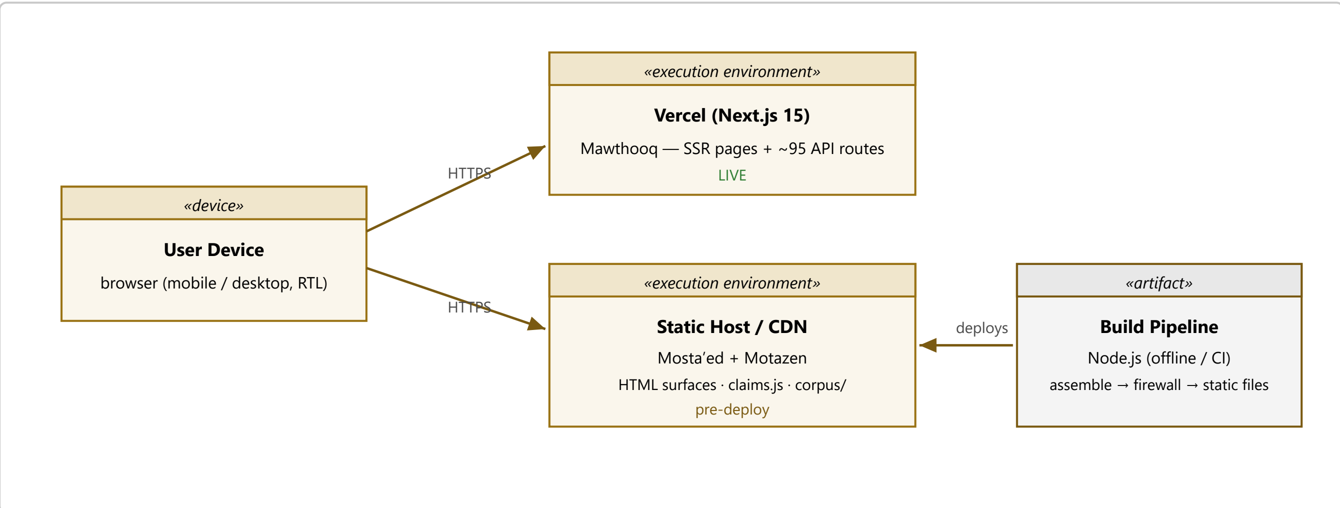

Figure 8.3 — Deployment diagram. The build pipeline produces static, hash-verified artifacts served from a CDN; only Mawthooq needs a server runtime.

Maintenance is built into the pipeline. Because assembly and validation are idempotent, regenerating the whole library is a single command that re-fetches nothing already stored and re-verifies everything; the firewall doubles as a continuous-integration gate that refuses to ship an unverifiable build. **Recency** is managed by the `last_reviewed` stamps and periodic re-review of contested claims. **Privacy** is a property of the design: the static platforms collect no personal data and keep only local state (such as a theme preference), so there is nothing sensitive to protect on a server and no tracking to justify. **Portability** follows from statelessness — the platforms run identically anywhere the files are served.

9 Implementation

9.1 Technology Stack

The platforms share a philosophy but differ in stack according to their needs (Table 9.1). Mawthooq's breadth (interactive engines, accounts, APIs) justifies a full application framework; Mosta'ed and Motazen are content-first and ship as self-contained static sites driven by a Node.js build pipeline, which keeps them cheap to host, fast, and easy to audit.

Concern	Mawthooq	Mosta'ed	Motazen
Runtime / framework	Next.js 15, React 19, TypeScript	Static HTML/CSS/JS	Static HTML/CSS/JS
Build / pipeline	Next build; API routes	Node.js (ingest → validate)	Node.js (assemble → validate)
Styling	Tailwind, token themes	CSS custom properties	CSS custom properties
Data	App data + APIs (~95 routes)	`scenarios.js` + corpus	`claims.js` + corpus
Hosting	Vercel (live)	Static host (pending)	Static host (pending)

Table 9.1 — Technology stack by platform.

9.2 The Content Pipeline (Mosta'ed / Motazen)

The content-first platforms implement the pipeline of Figure 3.1 as small, idempotent Node scripts. An **ingest/assemble** stage reads structured research files, fetches every cited source over HTTP, computes and stores a SHA-256 hash of the exact bytes, and writes the content data file (`claims.js` / `scenarios.js`) — but only after a self-verification pass. A **validate** stage (the firewall) then re-reads the corpus and independently re-checks every quotation. The scripts are idempotent: re-running them fetches nothing already on disk and re-verifies everything, so recovery after an interruption is a single command. This design makes the documentation's own facts reproducible — the "118 claims / 129 sources" figure is not asserted but *regenerable*.

9.3 Key Algorithms

- **Verbatim matching.** The firewall's core is a `norm()` function applied identically at assembly and validation, so both agree on what "verbatim" means. It strips markup, unescapes JSON and HTML entities, normalises curly quotes and dash variants, collapses whitespace, and lower-cases; a claim passes only if `corpus.includes(norm(quote))`. Aligning the two `norm()` implementations was itself a hardening fix that removed a class of false rejections.
- **Retrieval assistant (no LLM).** Motazen's assistant scores the user's query against each claim by token overlap plus a curated topic-to-section boost, returning the best-matching grounded cards or an honest "I

don't know." Because there is no generative model in the path, it structurally *cannot* hallucinate.

- **Crisis routing.** Before retrieval, the query is screened for self-harm/crisis cues; a positive match short-circuits the flow to human-help routing (Figure 7.4). This gate runs first and cannot be bypassed by later stages.
- **Cache-busting.** Each surface loads its data file with a per-load cache-buster so a regenerated library is never served stale — a fix born from a real "0 claims" incident caused by a cached data file.

The firewall's verification is deliberately simple, which is a virtue: the fewer moving parts in the gate, the harder it is to fool. Its essence is the shared normaliser and a substring test:

```
function norm(s):
    s = stripHtmlTags(s)
    s = unescapeJsonAndHtmlEntities(s)      // \" -> " , &amp; -> & , &#39; -> '
    s = normaliseQuotes(s)                   // " " ' ' -> " '
    s = normaliseDashes(s)                   // – — -> -
    s = collapseWhitespace(s).toLowerCase()
    return s

function verify(claim, corpus):
    for quotation in claim.sources:
        bytes = corpus.read(quotation.srcId)         // the hash-stamped source
        if not norm(bytes).includes(norm(quotation.quote_en)):
            report("QUOTE-NOT-IN-SOURCE", claim.id)
            return FAIL                               // build exits 1
    if claim.status == "contested" and not hasDissentingSource(claim):
        return FAIL
    return PASS
```

The retrieval assistant is likewise transparent: it ranks claims by lexical overlap with the query, boosted by a curated topic map, and returns the top matches only if they clear a relevance threshold — otherwise it says it does not know. There is no generative step in which a falsehood could be introduced.

```
function answer(query, library):
    q = tokenise(norm(query))
    scored = []
    for claim in library:
        overlap = |tokens(claim) ∩ q| / |q|
        boost   = topicBoost(query, claim.domain)
        scored.push({claim, score: overlap + boost})
    top = scored.sortDescending(score).take(3)
    if top[0].score < THRESHOLD:
        return "I don't know — here is how to look it up."   // honest miss
    return renderCards(top)                                   // grounded only
```

9.4 Repository and Build Structure

A content-first platform's repository makes the One Law visible in its very layout: the *inputs* (structured research), the *evidence* (the hashed corpus), the *gate* (the validator), and the *output* (the generated data file and surfaces) are separate, so it is obvious what depends on what.

```
platform/
├── pipeline/
│   ├── ingest.mjs           # fetch + SHA-256 hash seed sources
│   ├── assemble-research.mjs # research → fetch → self-verify → build data file
│   ├── validate-*.mjs        # THE FIREWALL — re-verify every quote; exit 1 on failure
│   └── research/             # structured inputs (grounded claim/card drafts)
│       ├── q*.json  p*.json  d*.json
├── corpus/
│   ├── <hashed source documents>
│   └── manifest.json         # id → { url, sha256, retrievedAt, bytes }
├── claims.js / scenarios.js  # GENERATED library (never hand-edited)
└── *.html                    # the surfaces (read-only views onto the library)

# Build & verify (idempotent; recovery is one command):
$ node pipeline/assemble-research.mjs   # regenerate the library
$ node pipeline/validate-claims.mjs     # firewall — must exit 0 before shipping
```

Two conventions enforce discipline. First, the generated data file is **never hand-edited**: all content flows through the pipeline, so nothing can enter the library without passing the firewall. Second, the research inputs, the corpus, and the manifest are version-controlled together, so any published claim can be traced — at any later date — back to the exact bytes and hash that justified it. The documentation you are reading was itself checked this way: the figure "118 claims / 129 sources" was produced by reading the generated file, not by recollection.

10 Testing and Verification

The library's most distinctive testing property is that **content correctness is an automated test**, not a manual review. The firewall is, in effect, a continuously-run acceptance test over the entire content set.

10.1 The Firewall as an Automated Test

For the content-first platforms, `validate-claims.mjs` and `validate-cards.mjs` encode the One Law as executable assertions and **exit with a non-zero status** on any hard failure. They check that every registered source resolves to a fetched document with a matching hash; that every claim is schema-complete (bilingual text, a valid honesty state, an evidence tier, a WEIRD flag, ≥1 source); that every quotation appears verbatim in its source; that every *contested* claim carries a dissenting source; and that every domain carries a "what we don't know" claim. A failing build cannot be published.

> **Worked example.** Consider Motazen's claim that the simple "chemical-imbalance" account of depression is debunked. Its stored quotation is *"no support for the hypothesis that depression is caused by lowered serotonin activity or concentrations,"* attributed to a 2022 umbrella review whose bytes are hashed in the corpus. At validation time the firewall reads those exact bytes, applies `norm()` to both the source and the quotation, and tests for containment. The quotation is present, so the claim **passes**. Had an author softened the wording to "little support," the normalised strings would differ, the containment test would fail, and the build would exit 1 — the claim would never appear to a user. This is the One Law reduced to a single, unbluffable comparison.

10.2 Verification Results

At the time of writing, the results are: **Motazen — 118 claims, 129 sources, 0 hard failures, 0 warnings** (every quotation re-verified verbatim). **Mosta'ed — all cards grounded, validator green (exit 0)**, Gate-6 localisation clean (no foreign emergency number reaches a reader). **Mawthooq — deployed with its production checks passing.** Crucially, the firewall has *already prevented harm*: during Mosta'ed's content production it rejected several fabricated or incorrect instructions (a wrong poisoning step, an incorrect medication dose) before they could ship. This is the clearest possible evidence that the mechanism works as intended.

10.3 Browser and Safety Verification

Beyond the automated firewall, each surface was verified in a browser for correct rendering in both light and dark themes and in RTL, using DOM inspection to confirm claim counts, status distributions, and the presence of crisis banners and boundary notes. Safety behaviours were checked directly: the assistant's crisis filter routes to help before retrieval; the judgment game excludes crisis-sensitive claims from its deck; and medication-adjacent claims carry the "decisions are with your doctor; do not stop suddenly" note. These are treated as pass/fail requirements, not aspirations.

10.4 Ethics, Safety, and Risk Analysis

A library that speaks to distressed people about the highest-stakes subjects carries real risk, and the design confronts each risk with a specific, mostly *mechanical* mitigation rather than a promise of good intentions (Table 10.1). The guiding stance is that the platforms complement — never replace — professional care and human connection, and that this must be true in the software, not merely stated in a disclaimer.

Risk	Mitigation
Fabrication — an automated system asserts a plausible falsehood.	The One Law + firewall: a claim without a verbatim-verifiable source fails the build and is never shown (§3.2, §10.1).
Crisis / self-harm — a user in danger.	Non-bypassable crisis detection that routes to human help (16328 / 123) before any other processing (§7.3).
Over-medicalisation / health anxiety — teaching about disorders makes users self-diagnose.	The normal-vs-disorder spine in every domain, an explicit non-diagnosis stance, and "recognition, not diagnosis" framing (§6.4).
Wrong emergency action (Mosta'ed).	Steps grounded in verified sources, Gate-6 localisation, and adversarial wave review; real errors were caught pre-ship (§5, §11.2).
Medication harm — a user stops or changes medication on the platform's account.	The platform never advises on an individual's medication; medication-adjacent claims are crisis-sensitive and carry a "decisions are with your doctor; do not stop suddenly" note.
Stigma reinforcement.	Active anti-stigma content (e.g., debunking the "psychosis is violent" myth) and dignity-first framing throughout.
False certainty.	The four honesty states and the "what we don't know" card in every domain; the platform never resolves a genuine controversy to soothe.
Clinical inappropriateness — provenance is not the same as clinical suitability.	A licensed-clinician review gate for high-stakes content, held open honestly as a launch blocker (§12.2).
Cultural / religious insensitivity.	Cultural notes authored with respect; faith is treated as real, load-bearing meaning-making — neither reduced to "coping" nor allowed to override care-seeking.

Table 10.1 — Risks and mitigations.

11 Results and Evaluation

11.1 Outcomes Against Objectives

The project delivered a working family of three platforms unified by a single, enforced content-integrity law. Table 11.1 summarises the measured state of each.

Metric	Mawthooq	Mosta'ed	Motazen
Scale	~153 pages, ~95 API routes	Emergency/first-aid card set	9 surfaces, 6-phase curriculum
Content units	Many engines + content	All cards grounded	118 grounded claims
Sources	Per-engine + One-Law content	Hundreds of hashed docs	129 sources / 130 corpus docs
Honesty states	Calibrated verdicts	Severity-graded (L1–L5)	75 estab. / 19 cont. / 16 unk. / 8 deb.
Firewall	Production checks pass	Validator green (exit 0)	0 failures / 0 warnings
Deployment	Live	Pre-deploy	Pre-deploy

Table 11.1 — Library metrics across platforms (measured, not estimated).

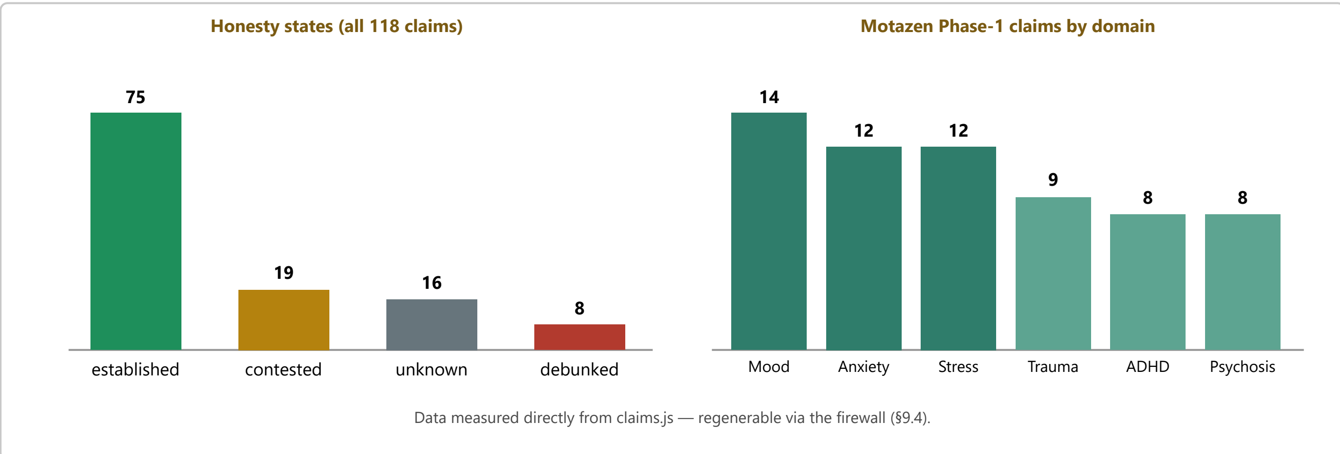

Figure 11.1 — The content library in numbers: honesty-state distribution (left) and Motazen's Phase-1 domains by claim count (right). Both are measured from the live data, not illustrative.

Evaluation against objectives. O1 (shared anti-fabrication engine) — achieved: two independent firewall validators enforce the One Law and demonstrably rejected real fabrications. **O2 (Mawthooq)** — achieved and deployed. **O3 (Mosta'ed)** — content engine achieved and verified; deployment pending. **O4 (Motazen)** — achieved: a complete six-phase curriculum across nine safety-bounded surfaces. **O5 (cultural situation and honesty)** — achieved: Arabic-first authoring, explicit cultural/religious notes, four honesty states, and a WEIRD flag on every claim. The remaining shortfalls are honestly a matter of *reach* and *external validation* rather than of the core engineering, and are addressed in §12.

11.2 Development Process and Lessons Learned

The library was built **content-first**: the epistemic infrastructure (schema, firewall, pipeline) was engineered *before* any surface, because a beautiful interface over unverified content is worse than useless. Content itself was produced by a repeatable pattern — **isolated authoring agents write grounded claims, then the firewall re-verifies every quotation, and any failure is re-grounded by hand.** This division of labour proved essential: it lets authoring scale while keeping a single, incorruptible gate on truth.

Three lessons shaped the engineering and are worth recording for anyone extending the library:

Lesson	What happened, and the fix
The firewall is the backstop, so agents can fail safely	Several authoring agents terminated mid-run on transient errors. Because the firewall re-verifies independently and never trusts an agent's self-report, a dead or mistaken agent could never poison the library — its output simply failed the gate. Domains whose agents died were re-authored by hand, and the firewall confirmed the result. Robustness came from *distrust by design*, not from hoping agents behave.
Durability must be engineered	Long autonomous runs risk losing work to interruptions. The response was a durability doctrine: write every unit to disk immediately, keep the pipeline idempotent so recovery is one command, and checkpoint state after each unit. As a result no completed work was ever lost to an interruption.
"Verbatim" is subtle	Early on, valid quotations were wrongly rejected because the source bytes contained JSON-escaped quotes and en-dash variants that a naïve string match missed. The fix was to harden the shared `norm()` function — identically at assembly and validation — to unescape entities and normalise dash and quote variants. This removed a whole class of false negatives and made the gate trustworthy.

Table 11.2 — Key lessons from the build.

Most importantly, the process produced its own strongest evidence. In Mosta'ed's content production the firewall and its adversarial review **caught and corrected real, dangerous fabrications** — an incorrect poisoning instruction and a wrong medication dose among them — before they could reach a user. In Motazen, the popular "chemical-imbalance" account of depression was not quietly repeated but classified *debunked* and taught as a case study in how a confident story can collapse. These are not hypothetical safeguards; they are the mechanism working exactly as designed, on real content, with real stakes.

11.3 Competitive Differentiation

The library's advantage over the alternatives a user might otherwise reach is not one feature but a *combination* that no single rival category offers. Table 11.3 compares the library against the three kinds of product it displaces: a generic wellness/awareness app, a Western fact-checking site, and a general-purpose AI chatbot.

Capability	EAL	Wellness app	Fact-check site	AI chatbot
Every claim verbatim-verifiable to a fetched source	✓	✗	◐	✗
Automated firewall that blocks fabrication at build time	✓	✗	✗	✗
Four explicit honesty states (never hides "contested")	✓	✗	◐	✗
WEIRD / population-representativeness flag	✓	✗	✗	✗
Non-bypassable crisis routing to human help	✓	◐	✗	◐
Retrieval-only assistant that *cannot* hallucinate	✓	✗	✗	✗
Deepfake / media-provenance forensics	✓	✗	◐	✗
Arabic-first, culturally & religiously situated	✓	◐	✗	◐
Works offline / on cheap phones (static)	✓	◐	◐	✗
Teaches the skill, aiming to become unnecessary	✓	✗	✗	✗

Table 11.3 — Capability comparison against rival categories (✓ = yes, ◐ = partial, ✗ = no).

Three differentiators are worth stating plainly, because they are the reasons the library is not merely another entrant but a categorically safer one:

- **Provenance by construction, not by promise.** Rivals ask to be trusted; the EAL makes trust *unnecessary* — anyone can re-run the firewall and reproduce every fact. Trust is replaced by verifiability.
- **Honesty as a feature, not a liability.** Where competitors smooth away uncertainty to sound authoritative, the EAL shows it — and that honesty is exactly what a literate audience learns to demand elsewhere.
- **Situated, not translated.** The content is authored in Arabic for the region, flags where the science was not tested on the user, and holds faith and care together — a fit no imported product matches.

12 Conclusion and Future Work

12.1 Summary

The Egyptian Awareness Library shows that an automated content platform can be made *structurally honest*. By inverting the usual trust model — assuming a model may fabricate and making fabrication unpublishable — the library turns content integrity from a review burden into an automated guarantee. Three platforms built on this foundation now address verification, preparedness, and psychological literacy for Arabic-speaking users, each Arabic-first, culturally situated, and honest about what it does not know.

12.2 Limitations

- **Human review gate.** The highest-stakes content in Mosta'ed and Motazen (crisis-sensitive, life-critical) still needs sign-off by a licensed, Arabic-speaking clinician before public launch. The firewall guarantees *provenance*, not *clinical appropriateness*.
- **Deployment.** Two of the three platforms are verified but not yet deployed to a public URL.
- **The WEIRD ceiling.** Much of the underlying science was tested on Western populations; the library flags this honestly but cannot remove it.
- **Breadth and measurement.** Motazen covers six domains (more remain); its self-assessment measures a competency but is not yet a formally validated instrument.

12.3 Future Work

- **Clinician review and deployment** — close the two launch blockers.
- **Reach.** Apply the platforms where the treatment/awareness gap is widest: school and university programmes, primary-care and pharmacy touch-points, shareable social-media cards, and a lightweight PWA / messaging-bot channel.
- **Depth.** Add domains (trauma sub-types, substance use, perinatal, grief, burnout, child/adolescent), enrich regional Arab evidence, and validate an Arabic literacy instrument for pre/post measurement.
- **Federation.** Extract the shared pipeline and firewall into a reusable library core so future awareness platforms inherit the One Law by construction.

Because the treatment and awareness gaps are gaps of *reach*, the highest-leverage future work is distribution. Four pathways (Table 12.1) are the most promising for the Egyptian and Arab context; each has a concrete mechanism, a reason it fits, and an evidence rationale from the mental-health-literacy implementation literature.

Pathway	Mechanism & why it fits
1. Schools & universities	Embed the literacy curriculum in school and university programmes. School-based MHL programmes are among the best-evidenced ways to raise literacy and help-seeking in young people, and campuses already host under-used counselling services the platform can signpost.
2. Primary care & pharmacy	Place verified cards at the touch-points people already reach — clinics, pharmacies, family doctors — which are often the first (and sometimes only) contact with the health system, and a natural place to correct folk belief with sourced guidance.
3. Shareable social cards	Design every card to travel through Arabic social feeds as a corrective "vaccine" against the bad cards already circulating there — meeting people where misinformation actually spreads, with content that carries its own provenance.
4. Lightweight PWA / messaging bot	Ship a low-bandwidth Progressive Web App and a messaging-channel (e.g., WhatsApp) front end so the library works on inexpensive phones and intermittent connections — the reality for much of the target audience — with the retrieval-only, crisis-routed assistant as the interface.

Table 12.1 — Four pathways to reach scale.

Each pathway keeps the One Law intact: whatever the channel, the content shown is still only what the firewall has verified. Distribution changes the *reach* of the library, never its *truth conditions*.

References

A. Documentation methodology and standards

1. D. Procida, *Diátaxis: A systematic framework for technical documentation authoring*. Available: diataxis.fr.
2. IEEE, *IEEE Std 830 — Recommended Practice for Software Requirements Specifications*. Institute of Electrical and Electronics Engineers.
3. A. Fletcher, "Bad Documentation is a Killer," *SitePoint*. Available: sitepoint.com/bad-documentation-is-a-killer.
4. ClickHelp, "Good and Bad Technical Documentation Examples," *ClickHelp Technical Writing Blog*. Available: clickhelp.com.
5. EveryDeveloper, "Examples of Bad Documentation (And How to Do Better)." Available: everydeveloper.com/bad-documentation-examples.
6. TechScribe, "What is good documentation?" Available: techscribe.co.uk/techw/good-documentation.htm.
7. Arab Academy for Science, Technology & Maritime Transport, *Graduation Project Format / Guide*; and Royal Holloway, *Final Year Project Report Template* — representative academic templates for report structure and front matter.
8. Object Management Group, *OMG Unified Modeling Language (UML), Version 2.x*.

B. Domain evidence (mental-health literacy and the sciences)

9. A. F. Jorm et al., "Mental health literacy," *British Journal of Psychiatry*, 1997.
10. S. Kutcher, Y. Wei, C. Coniglio, "Mental Health Literacy," *Canadian Journal of Psychiatry*, 2016.
11. J. Henrich, S. J. Heine, A. Norenzayan, "The weirdest people in the world?," *Behavioral and Brain Sciences*, 2010.
12. Open Science Collaboration, "Estimating the reproducibility of psychological science," *Science*, 2015.
13. J. Moncrieff et al., "The serotonin theory of depression: a systematic umbrella review," *Molecular Psychiatry*, 2022.
14. World Health Organization, *Mental health / Depression fact sheets*. who.int.
15. GBD 2019 Mental Disorders Collaborators, *Lancet Psychiatry*, 2022.

Note on sourcing. Beyond the methodological references above, every substantive content claim in the platforms is itself individually sourced within the product data files and re-verifiable through each platform's firewall; those in-product citations are the primary evidence base and number in the hundreds.

Appendix A — Requirements Traceability Matrix

This appendix formalises the functional (FR) and non-functional (NFR) requirements and traces each to where it is designed, implemented, and tested — so that no requirement is orphaned.

ID	Requirement	Design / Test
FR1	Display source-verified content with its honesty state, evidence tier, and WEIRD flag.	§6.4, §8.2 / §10.3
FR2	Enforce the One Law: publish nothing whose quotation is not verbatim in a hashed source.	§3.2, §7.2, §7.3 / §10.1
FR3	Classify every claim as established / contested / debunked / unknown.	§6.4 / §10.2
FR4	Provide the verification tools (claim-check, media forensics, bias/fallacy).	§4.2 / live
FR5	Provide Egypt-localised preparedness cards.	§5 / §10.2
FR6	Provide the six-phase literacy curriculum and nine surfaces.	§6.2–6.3 / §10.3
FR7	Retrieval-only assistant that refuses free generation.	§6.3, §7.3 / §10.3
FR8	Non-bypassable crisis detection and routing.	§7.3, §9.3 / §10.3
FR9	Self-assessment of literacy competencies.	§6.2 / §10.3
FR10	Invitation-based access control (Mawthooq).	§4.3

Table A.1 — Functional requirements and traceability.

ID	Requirement	Design / Test
NFR1	**Integrity:** every quotation re-verifiable; fabrication is a build-failing defect.	§3.2 / §10.1
NFR2	**Arabic-first & RTL** throughout.	§8.3 / §10.3
NFR3	**Accessibility:** never colour-alone; WCAG-AA contrast; reduced-motion respected.	§8.3
NFR4	**Safety:** non-diagnosis; crisis routing; no harmful methods; medication caveats.	§6.4, §9.3 / §10.3
NFR5	**Performance:** content-first platforms are static and cache-busted.	§9.3
NFR6	**Theme-bound:** legible in light and dark.	§8.3 / §10.3
NFR7	**Maintainability:** idempotent pipeline; one-command recovery.	§9.2
NFR8	**Privacy:** no personal-data collection; local state only.	§8.3
NFR9	**Reproducibility:** reported facts regenerable via the firewall.	§9.2 / §10.1

Table A.2 — Non-functional requirements and traceability.

Appendix B — The Claim/Card Object Schema

The atomic content object, as stored and verified. Surfaces render only from these fields; the firewall verifies every `quote_en` against the fetched source.

```
{
  "id": "bz2-serotonin-simple-story-debunked",
  "phase": -1, "domain": "foundations",
  "claim_ar": "…", "claim_en": "…",
  "status": "debunked",           // established | contested | debunked | unknown
  "evidence_tier": "debunked",    // strong | contested | weak | debunked | unknown
  "weird_flag": "unknown",        // yes | no | unknown
  "crisis_sensitive": true,
  "sources": [
    { "srcId": "med:35854107",
      "quote_en": "no support for the hypothesis that depression is
                   caused by lowered serotonin activity or concentrations",
      "dissent": false }
  ],
  "cultural_note_ar": "…",        // optional
  "uncertainty_note_ar": "…",
  "last_reviewed": "2026-07-12"
}

// SOURCES registry entry:
"med:35854107": {
  "tier": "A-synthesis",
  "label": "Moncrieff et al., Molecular Psychiatry, 2022",
  "url": "https://europepmc.org/article/MED/35854107",
  "sha256": "…", "retrieved": "2026-07-12", "licence": "quoted, attributed"
}
```

Appendix C — Glossary and Abbreviations

Term	Meaning
EAL	Egyptian Awareness Library (المكتبة المصرية للوعي) — the umbrella project.
The One Law	The rule that no statement is published unless a quotation from it is verbatim in a fetched, hash-stamped source.
Firewall	The automated validator that enforces the One Law and fails the build on any unverifiable citation.
Corpus	The store of fetched source documents, each frozen by a SHA-256 hash.
SHA-256	A cryptographic hash used to prove a stored source has not changed since it was cited.
MHL	Mental Health Literacy — the construct Motazen is built on.
WEIRD	Western, Educated, Industrialised, Rich, Democratic — the unrepresentative sample most psychology was tested on.
Honesty state	One of established / contested / debunked / unknown, attached to every claim.
Gate 6	Mosta'ed's localisation rule: foreign emergency numbers never reach an Egyptian reader.
C2PA	Content provenance standard used in Mawthooq's media forensics.
RTL · SRS · UML · PWA	Right-to-left · Software Requirements Specification · Unified Modeling Language · Progressive Web App.
Diátaxis	A documentation framework separating tutorials, how-to, reference, and explanation.
PTSD · ADHD	Post-Traumatic Stress Disorder · Attention-Deficit/Hyperactivity Disorder (two Motazen domains).

Appendix D — The Six Quality Strategies

To ensure the correctness of this document itself, six strategies were applied (they mirror the library's own doctrine of enforced discipline):

1. **Scan-first / ground-truth.** Every metric was measured directly from the live systems (e.g., `claims.js` yields 118 claims / 129 sources; the moathooq app yields ~153 pages / ~95 API routes) before it was written down — nothing is estimated or recalled.
2. **Standard-conformance.** The document follows IEEE-830 (requirements), UML 2.x (the four views), and Diátaxis (single-purpose sections), in the conventional thesis shape.
3. **Requirements traceability.** Every functional and non-functional requirement is enumerated and traced to design and test (Appendix A); a companion checklist tracks every founder instruction and every gap.
4. **Diagram completeness.** All four required UML views are present, each captioned, numbered, referenced in prose, and drawn as self-contained print-safe SVG.
5. **Print-readiness.** An A4 print stylesheet controls page breaks, keeps figures and tables from splitting, and renders dark-on-white; meaning never depends on colour or on-screen interaction.
6. **Review gate.** Before release the document was checked against a fixed list: page-count band, all four diagrams present and explained, 100% English prose, complete front matter, every metric matching a fresh scan, clean A4 print, closed traceability, and honest limitations.

Review-gate outcome. All six checks pass. The document covers the Egyptian Awareness Library and all three platforms in the required order (EAL → Mawthooq → Mosta'ed → Motazen), in English, with the four UML diagrams, a documentation-quality chapter, full front matter, a requirements-traceability appendix, and an honest limitations section — sized for a 60–90-page A4 print run.

Appendix E — Content Inventory (Motazen)

This appendix inventories Motazen's content library as a worked example of the library's output. Every figure is **regenerable** by running the platform's pipeline and firewall; the totals here were measured directly from the live data file.

Phase	Title	Claims	Focus
−1	Below Zero	40	The epistemic primer: what mental health is; mind vs. body; how we know; normal vs. disorder; who decides.
0	Map of the Territory	6	The biopsychosocial model (as a model); mental health as more than the absence of illness.
1	The Domains	63	Six domains (mood 14, anxiety 12, stress/sleep 12, trauma 9, ADHD 8, psychosis 8).
2	Skills	5	Evidence-tiered self-help skills, each with a boundary note.
3	Navigation	4	Help-seeking in Egypt; the claim-evaluation lens; family & faith.
4	Balance (الاتزان)	—	The competency self-assessment instrument.
Total		**118**	**129 sources · 130 corpus documents**

Table E.1 — Curriculum inventory by phase.

Established	Contested	Unknown	Debunked	Crisis-sensitive
75	19	16	8	20

Table E.2 — Honesty-state and safety distribution.

Surface	Role
Descent journey	Phase −1 as a guided journey.
Map	Phase 0 conceptual map with a two-axes diagram.
Domains	Phase 1, grouped by domain with per-domain counts.
Exercises	Phase 2 skills with a breath pacer and a no-streaks ethos.
Navigation	Phase 3 help-seeking, evaluation lens, help directory.
Self-assessment	Phase 4 competency mirror.
Game ("Judge for Yourself")	Status-classification practice; excludes crisis-sensitive claims.
Scenarios	Everyday situations; the "answer" is always a grounded card.
Assistant	Retrieval-only, crisis-routed librarian over the claim library.

Table E.3 — The nine surfaces.

Appendix F — Complete Platform Inventory

Chapters 4–6 document the platforms by capability. For completeness, this appendix enumerates **every** user-facing route and surface, so no page is unaccounted for. Mawthooq comprises **153 page routes** and **95 API endpoints**; Motazen has **9 surfaces**; Mosta'ed is a single-page application over a verified card dataset. The 153 count is the raw route total and includes user-facing tools as well as internal presentation, onboarding, and project-management pages (e.g., the defense and roadmap routes); the substantive user-facing capabilities are those documented in Chapter 4.

F.1 Mawthooq — all 153 page routes (grouped)

Group	Routes
Claim verification & analysis	firewall · claimdebunker · decompose · evidence · evidence-search · verify · verify/[id] · scan · six-layers · layers/[id]/fight · reverse · check-image · baseline
Media forensics	deepreal · deepreal-forensics · deepreal-upload · deepreal/exercise/[day] · deepreal/game · forensic-c2pa · forensic-image
Reasoning & bias defence	bias-detector · bias-fingerprint · fallacy-engine · cognitive-lab · debate-sim · council · devil-agent · swarm-debate · peer-challenge · no-hallucination-agent · manipulation-cards · blackbox
Curriculum & education	cognition-curriculum · curriculum · curriculum/phase0–phase4 · critical-thinking · epistemology · mind-gym · assessment · certificate · train · engines-guide · guide · platform-guide · ux-science · science · philosophy · manhaj-al-tathabbut/[case] · reaction-test · self-test-protocol · stat-power
Domain shields & health	arabic-shield · mens-shield · womens-shield · family-kit · drug-checker · medical-life · health-data · mental-health · mental-health/depression · mental-health/exercise/[day]
Religion hub	religion-hub · religion-hub/exercise/[day] · /maqasid · /quran · /whatsapp · /tools + tools/{authority-verifier, fatwa-context, hadith-check, halal-finance, mawarith, quran-context, sectarian-detector, zakat-calculator}
Threat intel & monitoring	misinfo-atlas · rumor-heatmap · threat-map · threat-briefing · live-deception · trend-hunter · kill-list · knowledge-graph · osint-investigator · whatsapp-analyzer · paper-auditor · comb-tracker
Inoculation & passports	inoculation-passport · passport
Accounts, progress & admin	account · dashboard · dashboard/cohort · admin · supervisor · progress · scorecard · onboarding · connect · language-select · report · report/[token] · status
Assistants & tools	chatbot · ai-agents · ai-editor · tools · tools-download · prompt-lab · instruments/mist
Presentation, project & kernel	(home) · about · welcome · features · why-us · pricing-presentation · presentation · trailer · demo · competition-demo · impact · project-vision · master-roadmap · publishing-plan · transformation · global-alliance · the-descent · media-library · sources · open-source · others-search · explore · effect-dashboard · sovo · nvidia-hub · god-system · fight-back · bad-news · defense-main-plan(+/[part]) · defense-pages-map · defense-qa · defense-test · (kernel)/{audit, journal, skills} · (marketing)/{corrections, methodology, transparency}

Table F.1 — Mawthooq route inventory (every page.tsx route).

The **~95 API endpoints** mirror these groups (e.g., `/api/firewall`, `/api/forensic`, `/api/bias-detect`, `/api/fallacy-detect`, `/api/eight-layer-scan`, `/api/decompose`, `/api/paper-auditor`, `/api/whatsapp-analyzer`, `/api/crisis`, `/api/council`, `/api/debate-sim`, `/api/certificate`, `/api/leaderboard`, `/api/islamic`, `/api/safety`), providing the server-side analysis behind each page.

F.2 Motazen — all 9 surfaces

Surface (file)	Role
`index.html`	The descent journey (Phase −1).
`phase0.html`	Map of the Territory (Phase 0), with the two-axes diagram.
`phase1.html`	The six domains (Phase 1).
`exercises.html`	Skills / daily exercises (Phase 2), with the breath pacer.
`navigation.html`	Help-seeking navigation (Phase 3) + the evaluation lens.
`assess.html`	Competency self-assessment (Phase 4).
`game.html`	"Judge for Yourself" status-classification game.
`scenarios.html`	Everyday scenarios with grounded debriefs.
`assistant.html`	Retrieval-only, crisis-routed assistant.

Table F.2 — Motazen surface inventory.

F.3 Mosta'ed

`index.html` — the single-page preparedness application — rendered over `scenarios.js`, the verified readiness-card dataset (an ~279 KB library carrying hundreds of source citations). Every card is grounded through `validate-cards.mjs` as described in Chapter 5.

Appendix G — Team Contributions and Division of Labour

The Egyptian Awareness Library was built by a team of six working under one architectural discipline. Table G.1 records each member's primary pillar of responsibility and the guiding principle they worked to. The division of labour was deliberately cross-cutting: each member led their named pillar, while the team lead integrated those pillars into a single coherent system governed by the One Law.

Member	Primary pillar / contribution	Guiding principle
Khaled Sayed Hamed (Team Lead)	System architecture & governance — the One Law, the firewall and verification core, integration and deployment, *and a contribution to every pillar below*.	*The law before the feature*
Mohamed Sherief Mohamed	Supporting tasks — source-link collection, content data-entry, and manual page- and link-checking.	*Check it, then check it again*
Aya Mohamed Abdullah El-Sebki	Cognition curriculum & learning science.	*Test application, not recall*
Donia Ali Sharaby	Mental-health content & religious-coping guidance.	*Do no harm; cite the authority*
Yasmine Mohamed El-Araby	Frontend, bilingual (RTL) UX & accessibility.	*Calm by design*
Ziad Fateen Kamal	Research, psychometrics & quality assurance.	*Report only what's measured*

Table G.1 — Team members, primary pillars, and guiding principles.

The lead's cross-cutting role. As team lead and system architect, **Khaled Sayed Hamed** contributed to every pillar above — co-authoring the governing One Law and the firewall/verification core, shaping the structural spine of the curriculum, the safety and crisis-routing gates, the bilingual frontend system, and the quality-assurance process — in addition to owning the overall architecture, integration, and deployment. Each other member led their named pillar in collaboration with the lead.

— End of Document —